

# Parallelism I: Inside the Core

# The final

- Comprehensive
- Same general format as the Midterm.
- Review the homeworks, the slides, and the quizzes.

# Key Points

- What is wide issue mean?
- How does does it affect performance?
- How does it affect pipeline design?
- What is the basic idea behind out-of-order execution?
- What is the difference between a true and false dependence?
- How do OOO processors remove false dependences?
- What is Simultaneous Multithreading?

# Parallelism

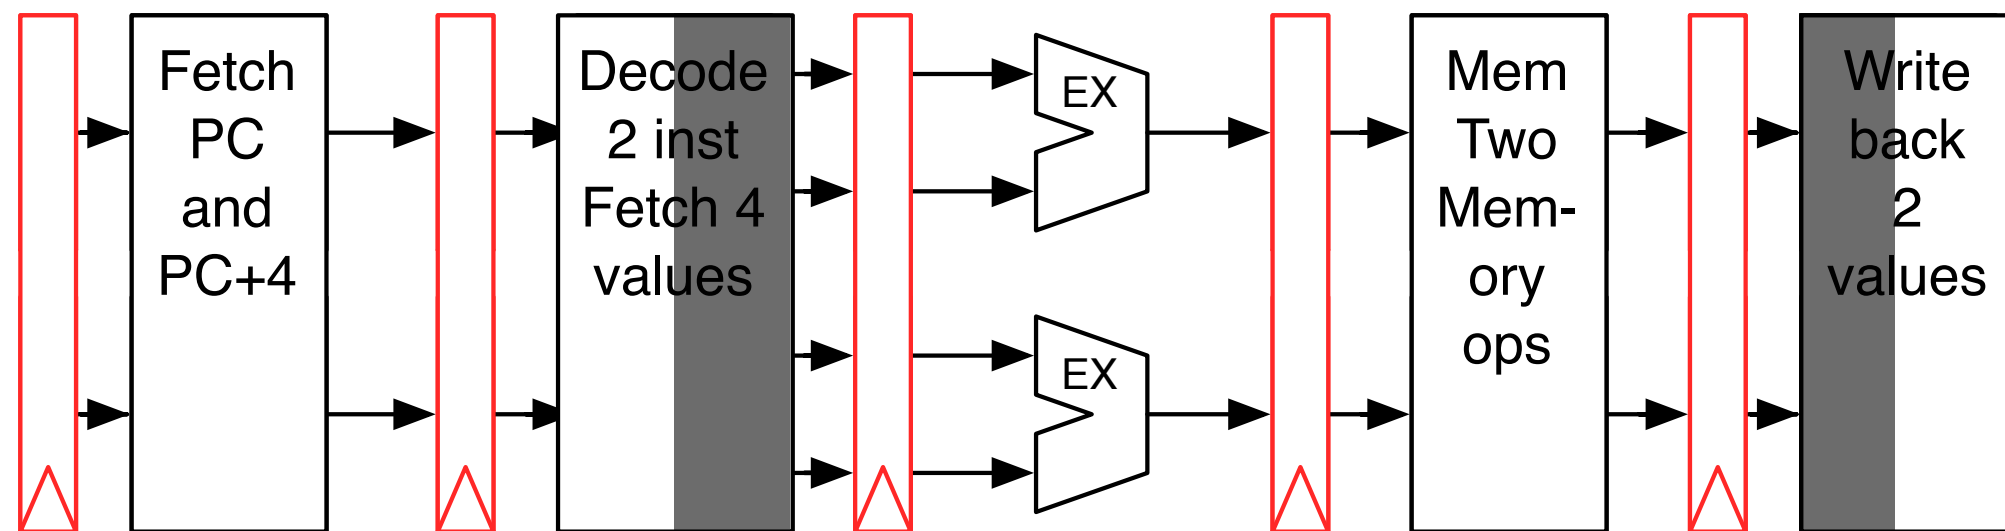
- $ET = IC * CPI * CT$
- IC is more or less fixed
- We have shrunk cycle time as far as we can
- We have achieved a CPI of 1.
- Can we get faster?

# Parallelism

- $ET = IC * CPI * CT$
- IC is more or less fixed
- We have shrunk cycle time as far as we can
- We have achieved a CPI of 1.
- Can we get faster?

We can reduce our CPI to less than 1.  
The processor must do multiple operations at once.  
This is called Instruction Level Parallelism (ILP)

# Approach 1: Widen the pipeline



- Process two instructions at once instead of 1
- Often 1 “odd” PC instruction and 1 “even” PC
  - This keeps the instruction fetch logic simpler.
- 2-wide, in-order, superscalar processor
- Potential problems?

# Single issue refresher

	cycle 0	cycle 1	cycle 2	cycle 3	cycle 4	cycle 5	cycle 6	cycle 7	cycle 8
add \$s1,\$s2,\$s3	F	D	E	M	W				
sub \$s2,\$s4,\$s5		F	D	E	M	W			
ld \$s3, 0(\$s2)			F	D	E	M	W		
add \$t1, \$s3, \$s3				F	D	D	E	M	W

# Single issue refresher

	cycle 0	cycle 1	cycle 2	cycle 3	cycle 4	cycle 5	cycle 6	cycle 7	cycle 8
add \$s1,\$s2,\$s3	F	D	E	M	W				
sub \$s2,\$s4,\$s5		F	D	E	M	W			
ld \$s3, 0(\$s2)			F	D	E	M	W		
add \$t1, \$s3, \$s3				F	D	D	E	M	W

Forwarding





# Single issue refresher

	cycle 0	cycle 1	cycle 2	cycle 3	cycle 4	cycle 5	cycle 6	cycle 7	cycle 8
add \$s1,\$s2,\$s3	F	D	E	M	W				
sub \$s2,\$s4,\$s5		F	D	E	M	W			
ld \$s3, 0(\$s2)			F	D	E	M	W		
add \$t1, \$s3, \$s3				F	D	D	E	M	W

Forwarding

Forwarding

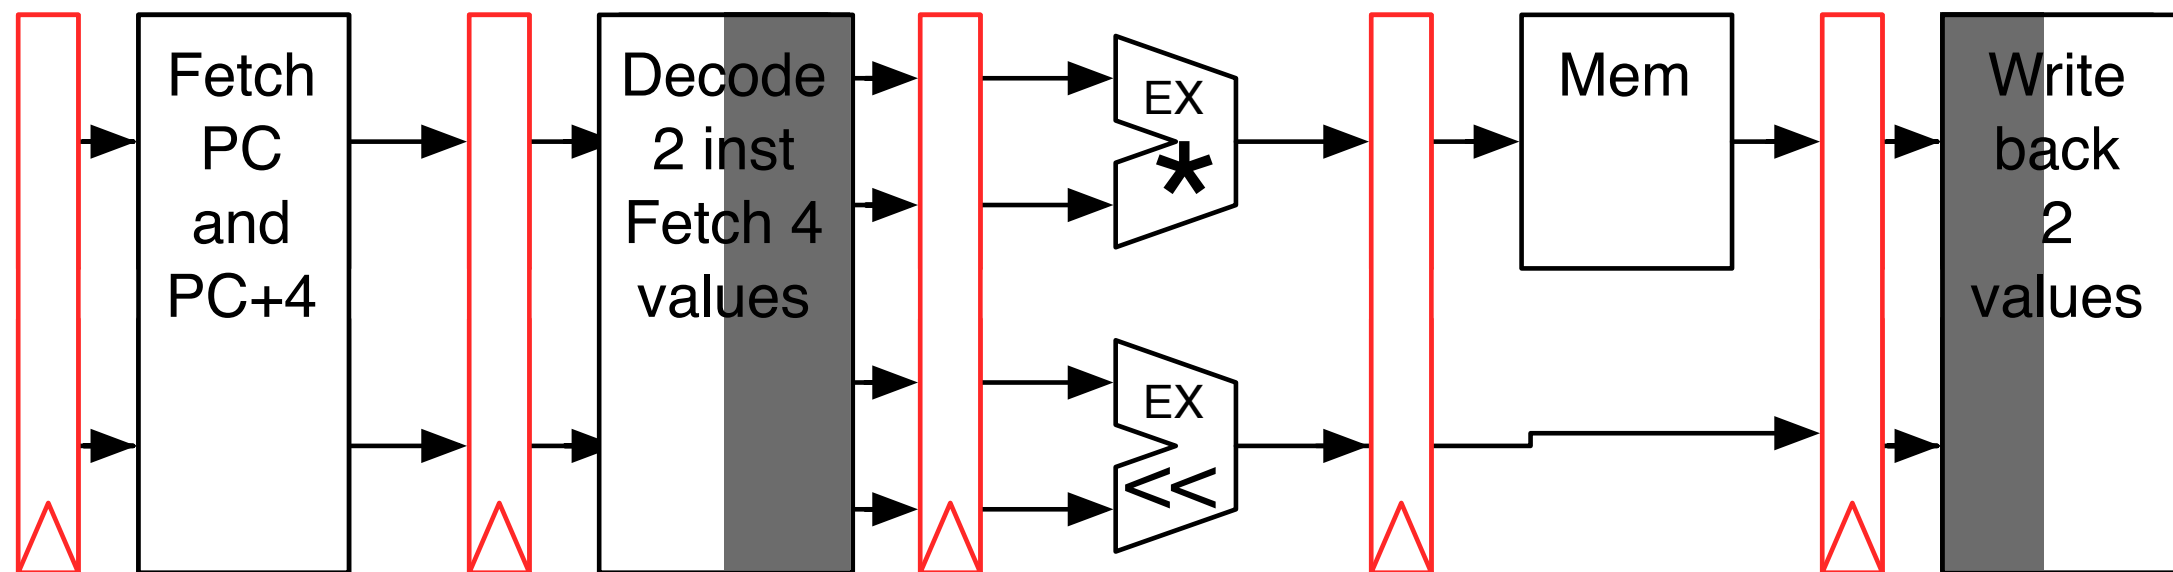
# Dual issue: Ideal Case

	0	1	2	3	4	5	6	7	8
add \$s1,\$s2,\$s3	F	D	E	M	W				
sub \$s2,\$s4,\$s5	F	D	E	M	W				
ld \$s3, 0(\$s2)		F	D	E	M	W			
add \$t1, \$s3, \$s3		F	D	E	M	W			
...			F	D	E	M	W		
...			F	D	E	M	W		
...				F	D	E	M	W	
...				F	D	E	M	W	
...					F	D	E	M	W
...					F	D	E	M	W

CPI == 0.5!

# Dual issue: Structural Hazards

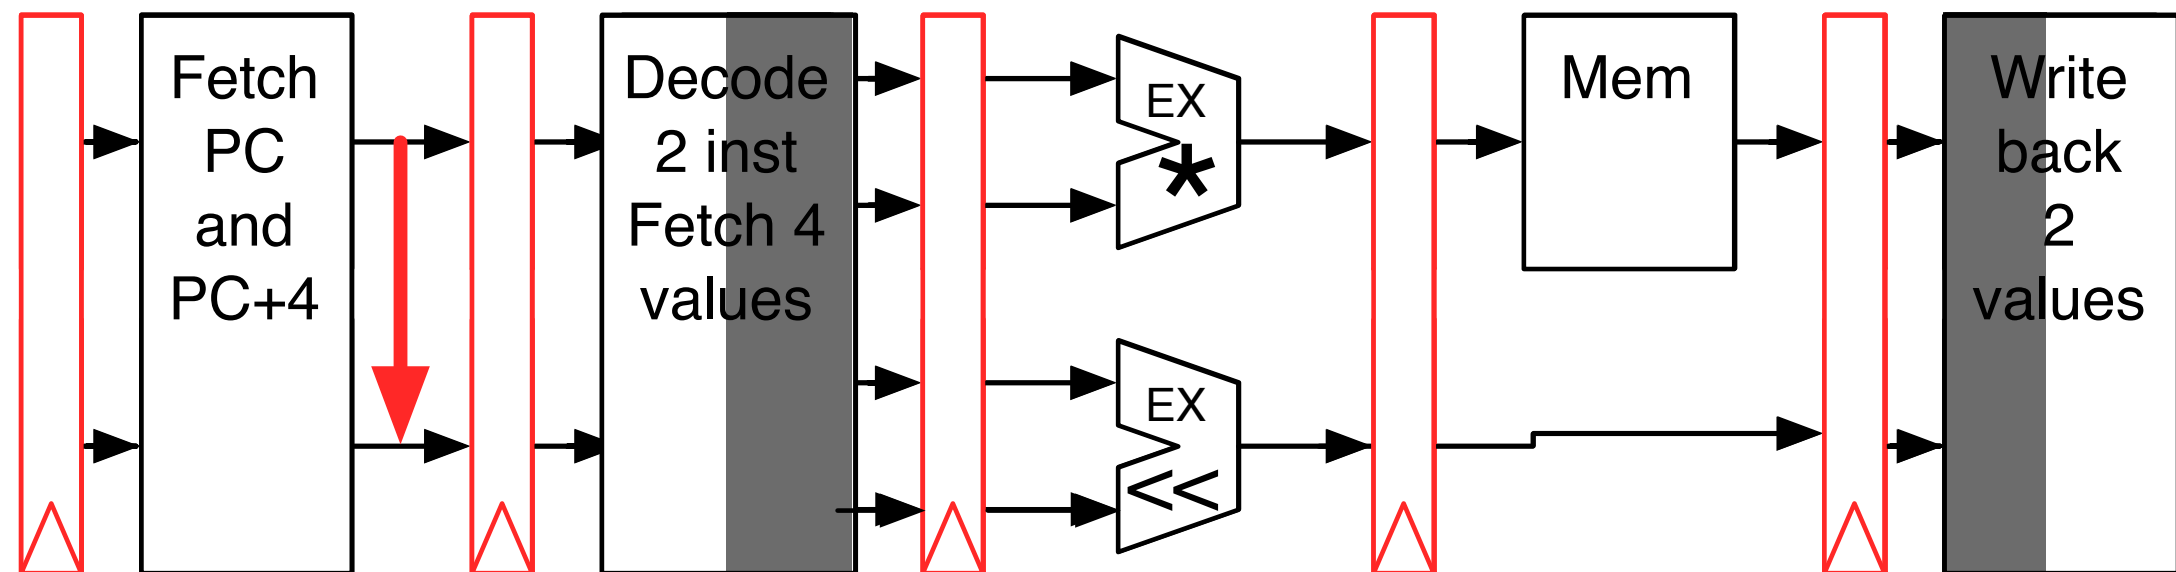
- Structural hazards
  - We might not replicate everything
  - Perhaps only one multiplier, one shifter, and one load/store unit
  - What if the instruction is in the wrong place?



If an “upper” instruction needs the “lower” pipeline,  
squash the “lower” instruction

# Dual issue: Structural Hazards

- Structural hazards
  - We might not replicate everything
  - Perhaps only one multiplier, one shifter, and one load/store unit
  - What if the instruction is in the wrong place?



If an “upper” instruction needs the “lower” pipeline,  
squash the “lower” instruction

# Dual issue: dealing with hazards

	0	1	2	3	4	5	6	7	8	9	10
add	F	D	E	M	W						
sub	F	D	E	M	W						
Mul		F	D	E	M	W					
Shift		F	D	E	M	W					
Shift			F	D	E	M	W				
Ld			F	x	x	x	x				
Shift				x	x	x	x	x			
Ld				F	D	E	M	W			

# Dual issue: dealing with hazards

PC = 0

	0	1	2	3	4	5	6	7	8	9	10
add	F	D	E	M	W						
sub	F	D	E	M	W						
Mul		F	D	E	M	W					
Shift		F	D	E	M	W					
Shift			F	D	E	M	W				
Ld			F	x	x	x	x				
Shift				x	x	x	x	x			
Ld				F	D	E	M	W			

# Dual issue: dealing with hazards

PC = 0		0	1	2	3	4	5	6	7	8	9	10
	add	F	D	E	M	W						
	sub	F	D	E	M	W						
PC = 8	Mul		F	D	E	M	W					
	Shift		F	D	E	M	W					
	Shift			F	D	E	M	W				
	Ld			F	x	x	x	x				
	Shift				x	x	x	x	x			
	Ld				F	D	E	M	W			

# Dual issue: dealing with hazards

		0	1	2	3	4	5	6	7	8	9	10
PC = 0	add	F	D	E	M	W						
	sub	F	D	E	M	W						
PC = 8	Mul		F	D	E	M	W					
	Shift		F	D	E	M	W					
PC = 12	Shift			F	D	E	M	W	Shift moves to lower pipe load is squashed			
	Ld			F	x	x	x	x				
	Shift				x	x	x	x	x			
	Ld				F	D	E	M	W			

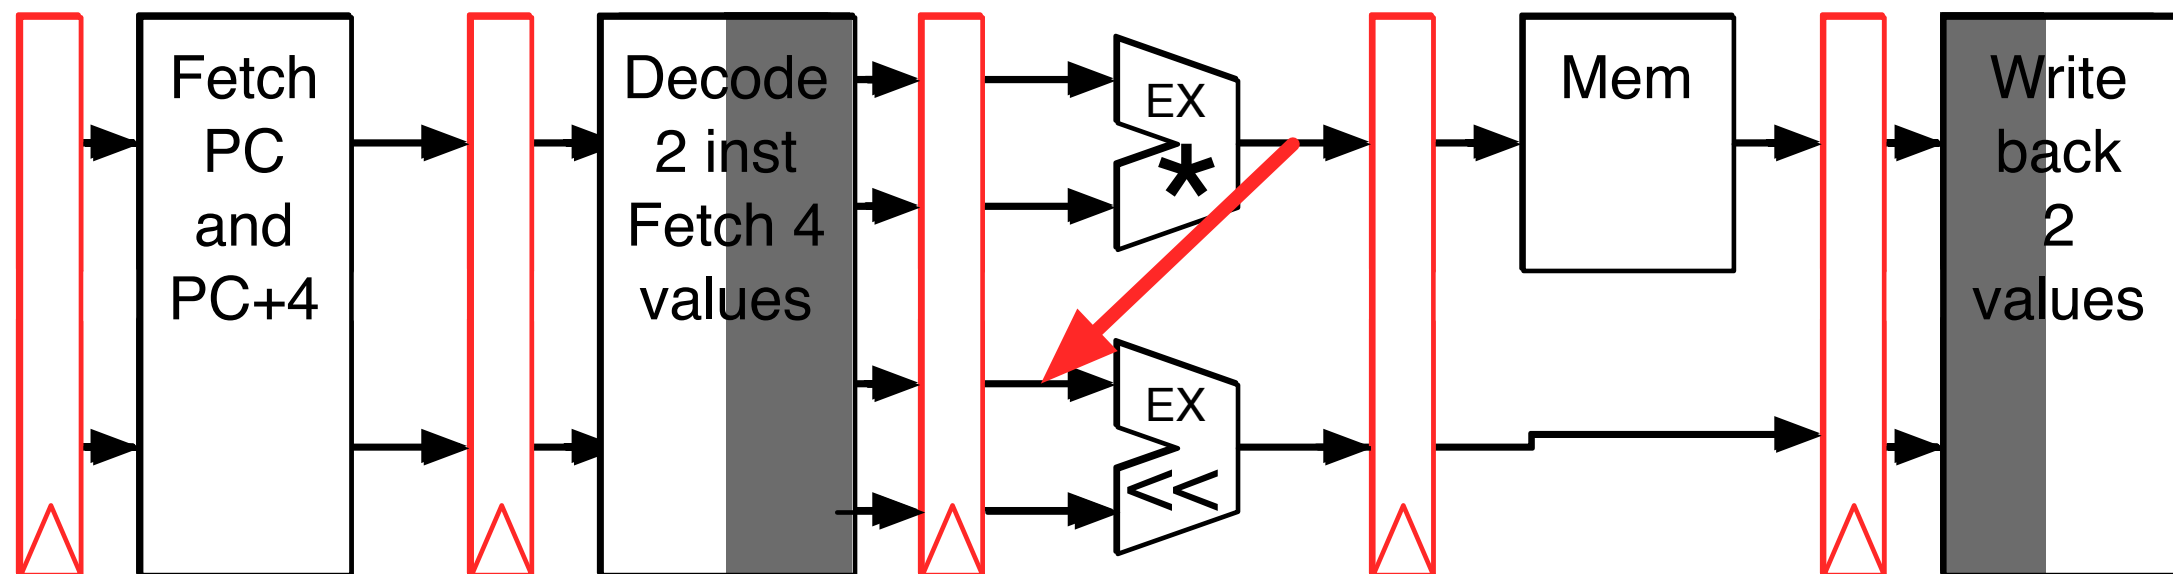


# Dual issue: dealing with hazards

		0	1	2	3	4	5	6	7	8	9	10
PC = 0	add	F	D	E	M	W						
	sub	F	D	E	M	W						
PC = 8	Mul		F	D	E	M	W					
	Shift		F	D	E	M	W					
PC = 12	Shift			F	D	E	M	W	Shift moves to lower pipe load is squashed			
	Ld			F	x	x	x	x				
PC = 12	Shift				x	x	x	x	x	Load uses lower pipe Shift becomes a noop		
	Ld				F	D	E	M	W			

# Dual issue: Data Hazards

- The “lower” instruction may need a value produced by the “upper” instruction
- Forwarding cannot help us -- we must stall.



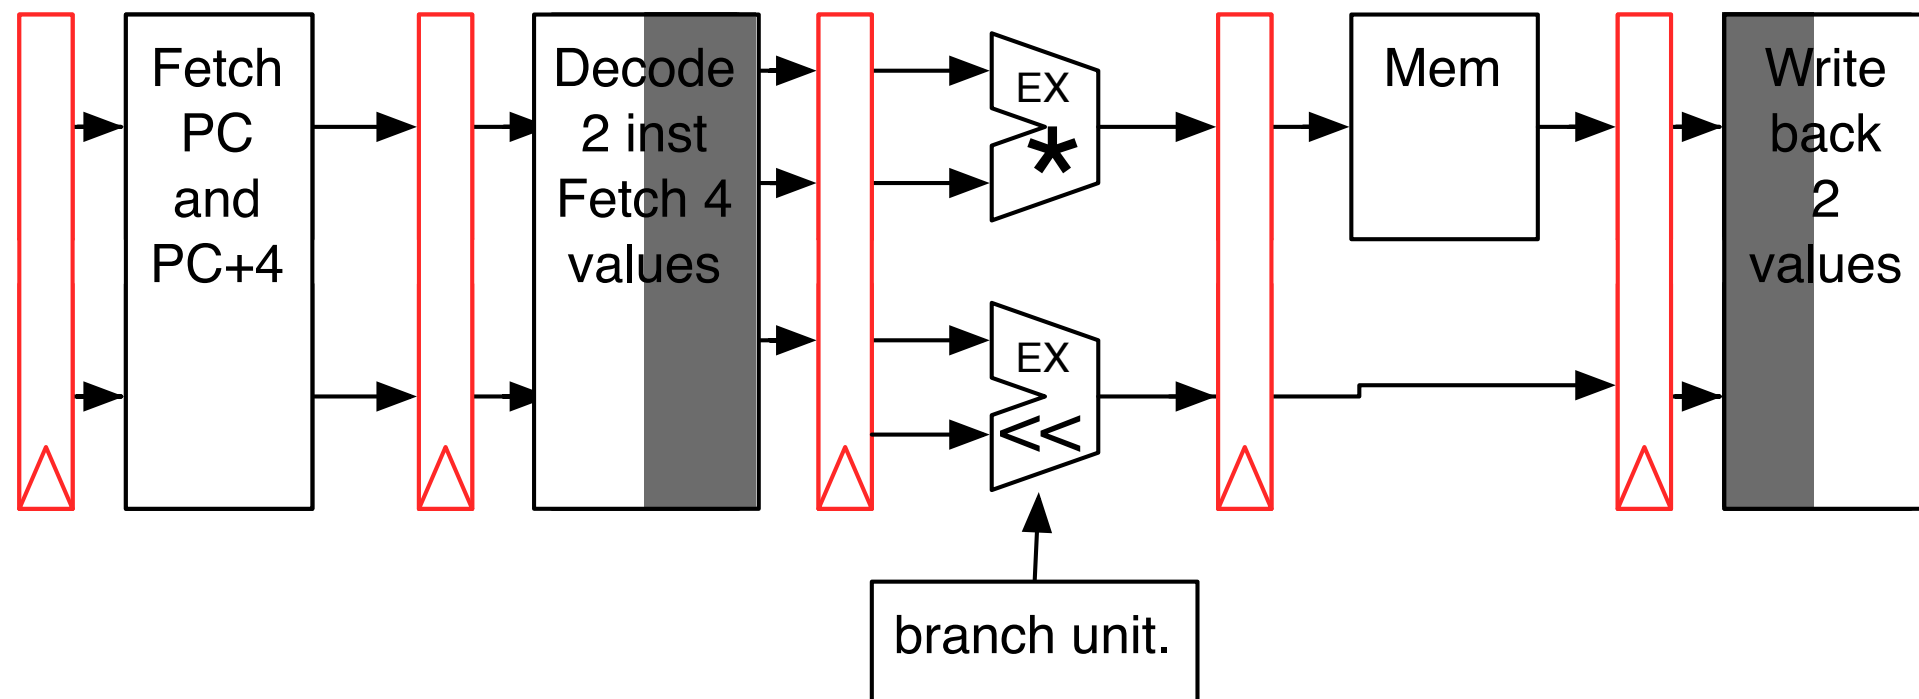
# Dual issue: dealing with hazards

- Forwarding is essential!
- Both pipes stall.

	0	1	2	3	4	5	6	7	8	9	10
add \$s1, \$s3, #4	F	D	E	M	W						
sub \$s4, \$s1, #4	F	D	D	E	M	W					
add ...		F	F	D	E	M	W				
sub ...		F	F	D	E	M	W				
and ...				F	D	E	M	W			
or ...				F	D	E	M	W			

# Dual issue: Control Hazards

- The “upper” instruction might be branch.
- The “lower” instruction might be on the wrong path
- Solution 1: Require branches to execute in the lower pipeline -- See “structural hazards”.
- What about consecutive branches? -- Exercise for the reader
- What about branches to odd addresses? -- Squash the upper pipe



# Beyond Dual Issue

- Wider pipelines are possible.
  - There is often a separate floating point pipeline.
- Wide issue leads to hardware complexity
- Compiling gets harder, too.
- In practice, processors use of two options if they want more ILP
  - Change the ISA and build a smart compiler: VLIW
  - Keep the same ISA and build a smart processors: Out-of-order


# Going Out of Order: Data dependence refresher.

1: add \$t1, \$s2, \$s3

2: sub \$t2, \$s3, \$s4


3: or \$t5, \$t1, \$t2

4: add \$t3, \$t1, \$t2



# Going Out of Order: Data dependence refresher.




```
1: add $t1, $s2, $s3
2: sub $t2, $s3, $s4
3: or  $t5, $t1, $t2
4: add $t3, $t1, $t2
```



**There is parallelism!!**  
**We can execute**  
**1 & 2 at once**  
**and**  
**3 & 4 at once**

# Going Out of Order: Data dependence refresher.

```
1: add $t1, $s2, $s3
2: sub $t2, $s3, $s4
3: or  $t5, $t1, $t2
4: add $t3, $t1, $t2
```



**There is parallelism!!**  
We can execute  
1 & 2 at once  
and  
3 & 4 at once

**We can parallelize instructions  
that do not have a “read-after-  
write” dependence (RAW)**



# Data dependences


- In general, if there is no dependence between two instructions, we can execute them in either order or simultaneously.

- But beware:

- Is there a dependence here?

1: add \$t1, \$s2, \$s3

2: sub \$t1, \$s3, \$s4



- Can we reorder the instructions?

2: sub \$t1, \$s3, \$s4

1: add \$t1, \$s2, \$s3

- Is the result the same?

# Data dependences


- In general, if there is no dependence between two instructions, we can execute them in either order or simultaneously.

- But beware:

- Is there a dependence here?

1: add \$t1, \$s2, \$s3

2: sub \$t1, \$s3, \$s4



- Can we reorder the instructions?

2: sub \$t1, \$s3, \$s4

1: add \$t1, \$s2, \$s3

**No! The final value of \$t1 is different**

- Is the result the same?

# False Dependence #1

- Also called “Write-after-Write” dependences (WAW) occur when two instructions write to the same value
- The dependence is “false” because no data flows between the instructions -- They just produce an output with the same name.

# Beware again!

- Is there a dependence here?

1: add \$t1, \$s2, \$s3

2: sub \$s2, \$s3, \$s4

1

2

- Can we reorder the instructions?

2: sub \$s2, \$s3, \$s4

1: add \$t1, \$s2, \$s3


- Is the result the same?

# Beware again!

- Is there a dependence here?

1: add \$t1, \$s2, \$s3

2: sub \$s2, \$s3, \$s4



- Can we reorder the instructions?

2: sub \$s2, \$s3, \$s4

1: add \$t1, \$s2, \$s3

**No! The value in \$s2 that I  
needs will be destroyed**

- Is the result the same?

# False Dependence #2

- This is a Write-after-Read (WAR) dependence
- Again, it is “false” because no data flows between the instructions

# Out-of-Order Execution

- Any sequence of instructions has set of RAW, WAW, and WAR dependences that constrain its execution.
- Can we design a processor that extracts as much parallelism as possible, while still respecting these dependences?

# The Central OOO Idea

1. Fetch a bunch of instructions
2. Build the dependence graph
3. Find all instructions with no unmet dependences
4. Execute them.
5. Repeat



# Example

1: add \$t1,\$s2,\$s3

2: sub \$t2,\$s3,\$s4

3: or \$t3,\$t1,\$t2

4: add \$t5,\$t1,\$t2

— WAR —→

— WAW —→

— RAW —→


# Example

1: add \$t1,\$s2,\$s3

2: sub \$t2,\$s3,\$s4

3: or \$t3,\$t1,\$t2

4: add \$t5,\$t1,\$t2



— WAR —→

— WAW —→

— RAW —→


# Example

1: add \$t1,\$s2,\$s3

2: sub \$t2,\$s3,\$s4

3: or \$t3,\$t1,\$t2

4: add \$t5,\$t1,\$t2



— WAR —→

— WAW —→

— RAW —→

# Example

1: add \$t1,\$s2,\$s3

2: sub \$t2,\$s3,\$s4

3: or \$t3,\$t1,\$t2


4: add \$t5,\$t1,\$t2

5: or \$t4,\$s1,\$s3

6: mul \$t2,\$t3,\$s5

7: sl \$t3,\$t4,\$t2

8: add \$t3,\$t5,\$t1




— WAR —→

— WAW —→

— RAW —→


# Example

```
1: add $t1,$s2,$s3
2: sub $t2,$s3,$s4
3: or  $t3,$t1,$t2
4: add $t5,$t1,$t2
5: or  $t4,$s1,$s3
6: mul $t2,$t3,$s5
7: sl  $t3,$t4,$t2
8: add $t3,$t5,$t1
```




# Example

```
1: add $t1,$s2,$s3
2: sub $t2,$s3,$s4
3: or  $t3,$t1,$t2
4: add $t5,$t1,$t2
5: or  $t4,$s1,$s3
6: mul $t2,$t3,$s5
7: sl  $t3,$t4,$t2
8: add $t3,$t5,$t1
```



# Example

```
1: add $t1,$s2,$s3
2: sub $t2,$s3,$s4
3: or  $t3,$t1,$t2
4: add $t5,$t1,$t2
5: or  $t4,$s1,$s3
6: mul $t2,$t3,$s5
7: sl  $t3,$t4,$t2
8: add $t3,$t5,$t1
```



# Example

```
1: add $t1, $s2, $s3
```

```
2:  sub    $t2, $s3, $s4
```

3: or  $\$t3, \$t1, \$t2$


```
4:  add  $t5, $t1, $t2
```

5: or  $\$t4, \$s1, \$s3$

```
6: mul    $t2, $t3, $s5
```

7: s1 \$t3, \$t4, \$t2

```
8:  add  $t3,$t5,$t1
```

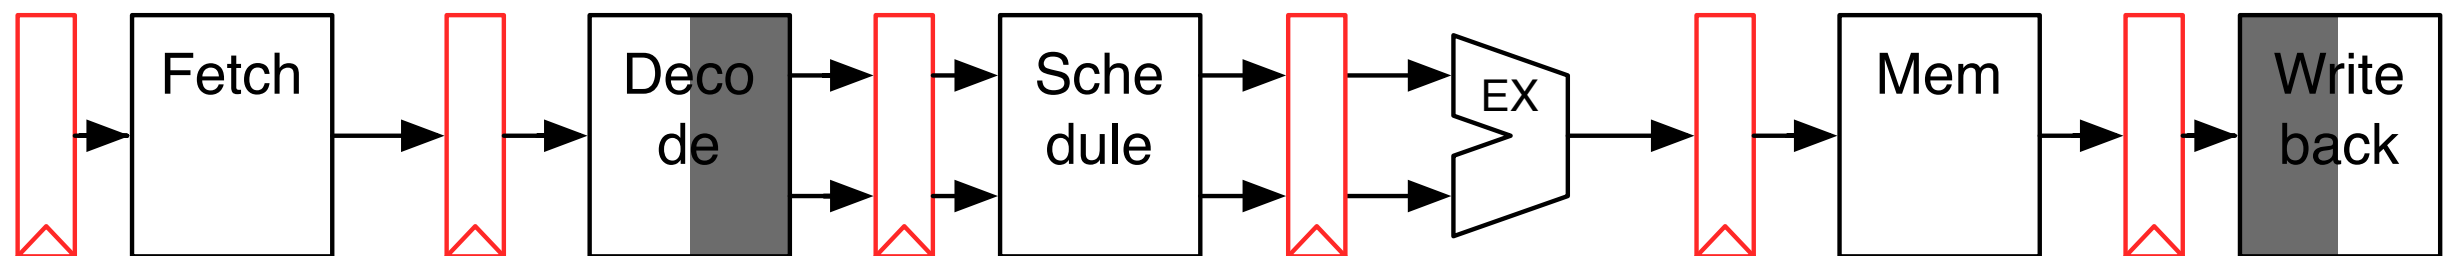


# 8 Instructions in 5 cycles



# Simplified OOO Pipeline

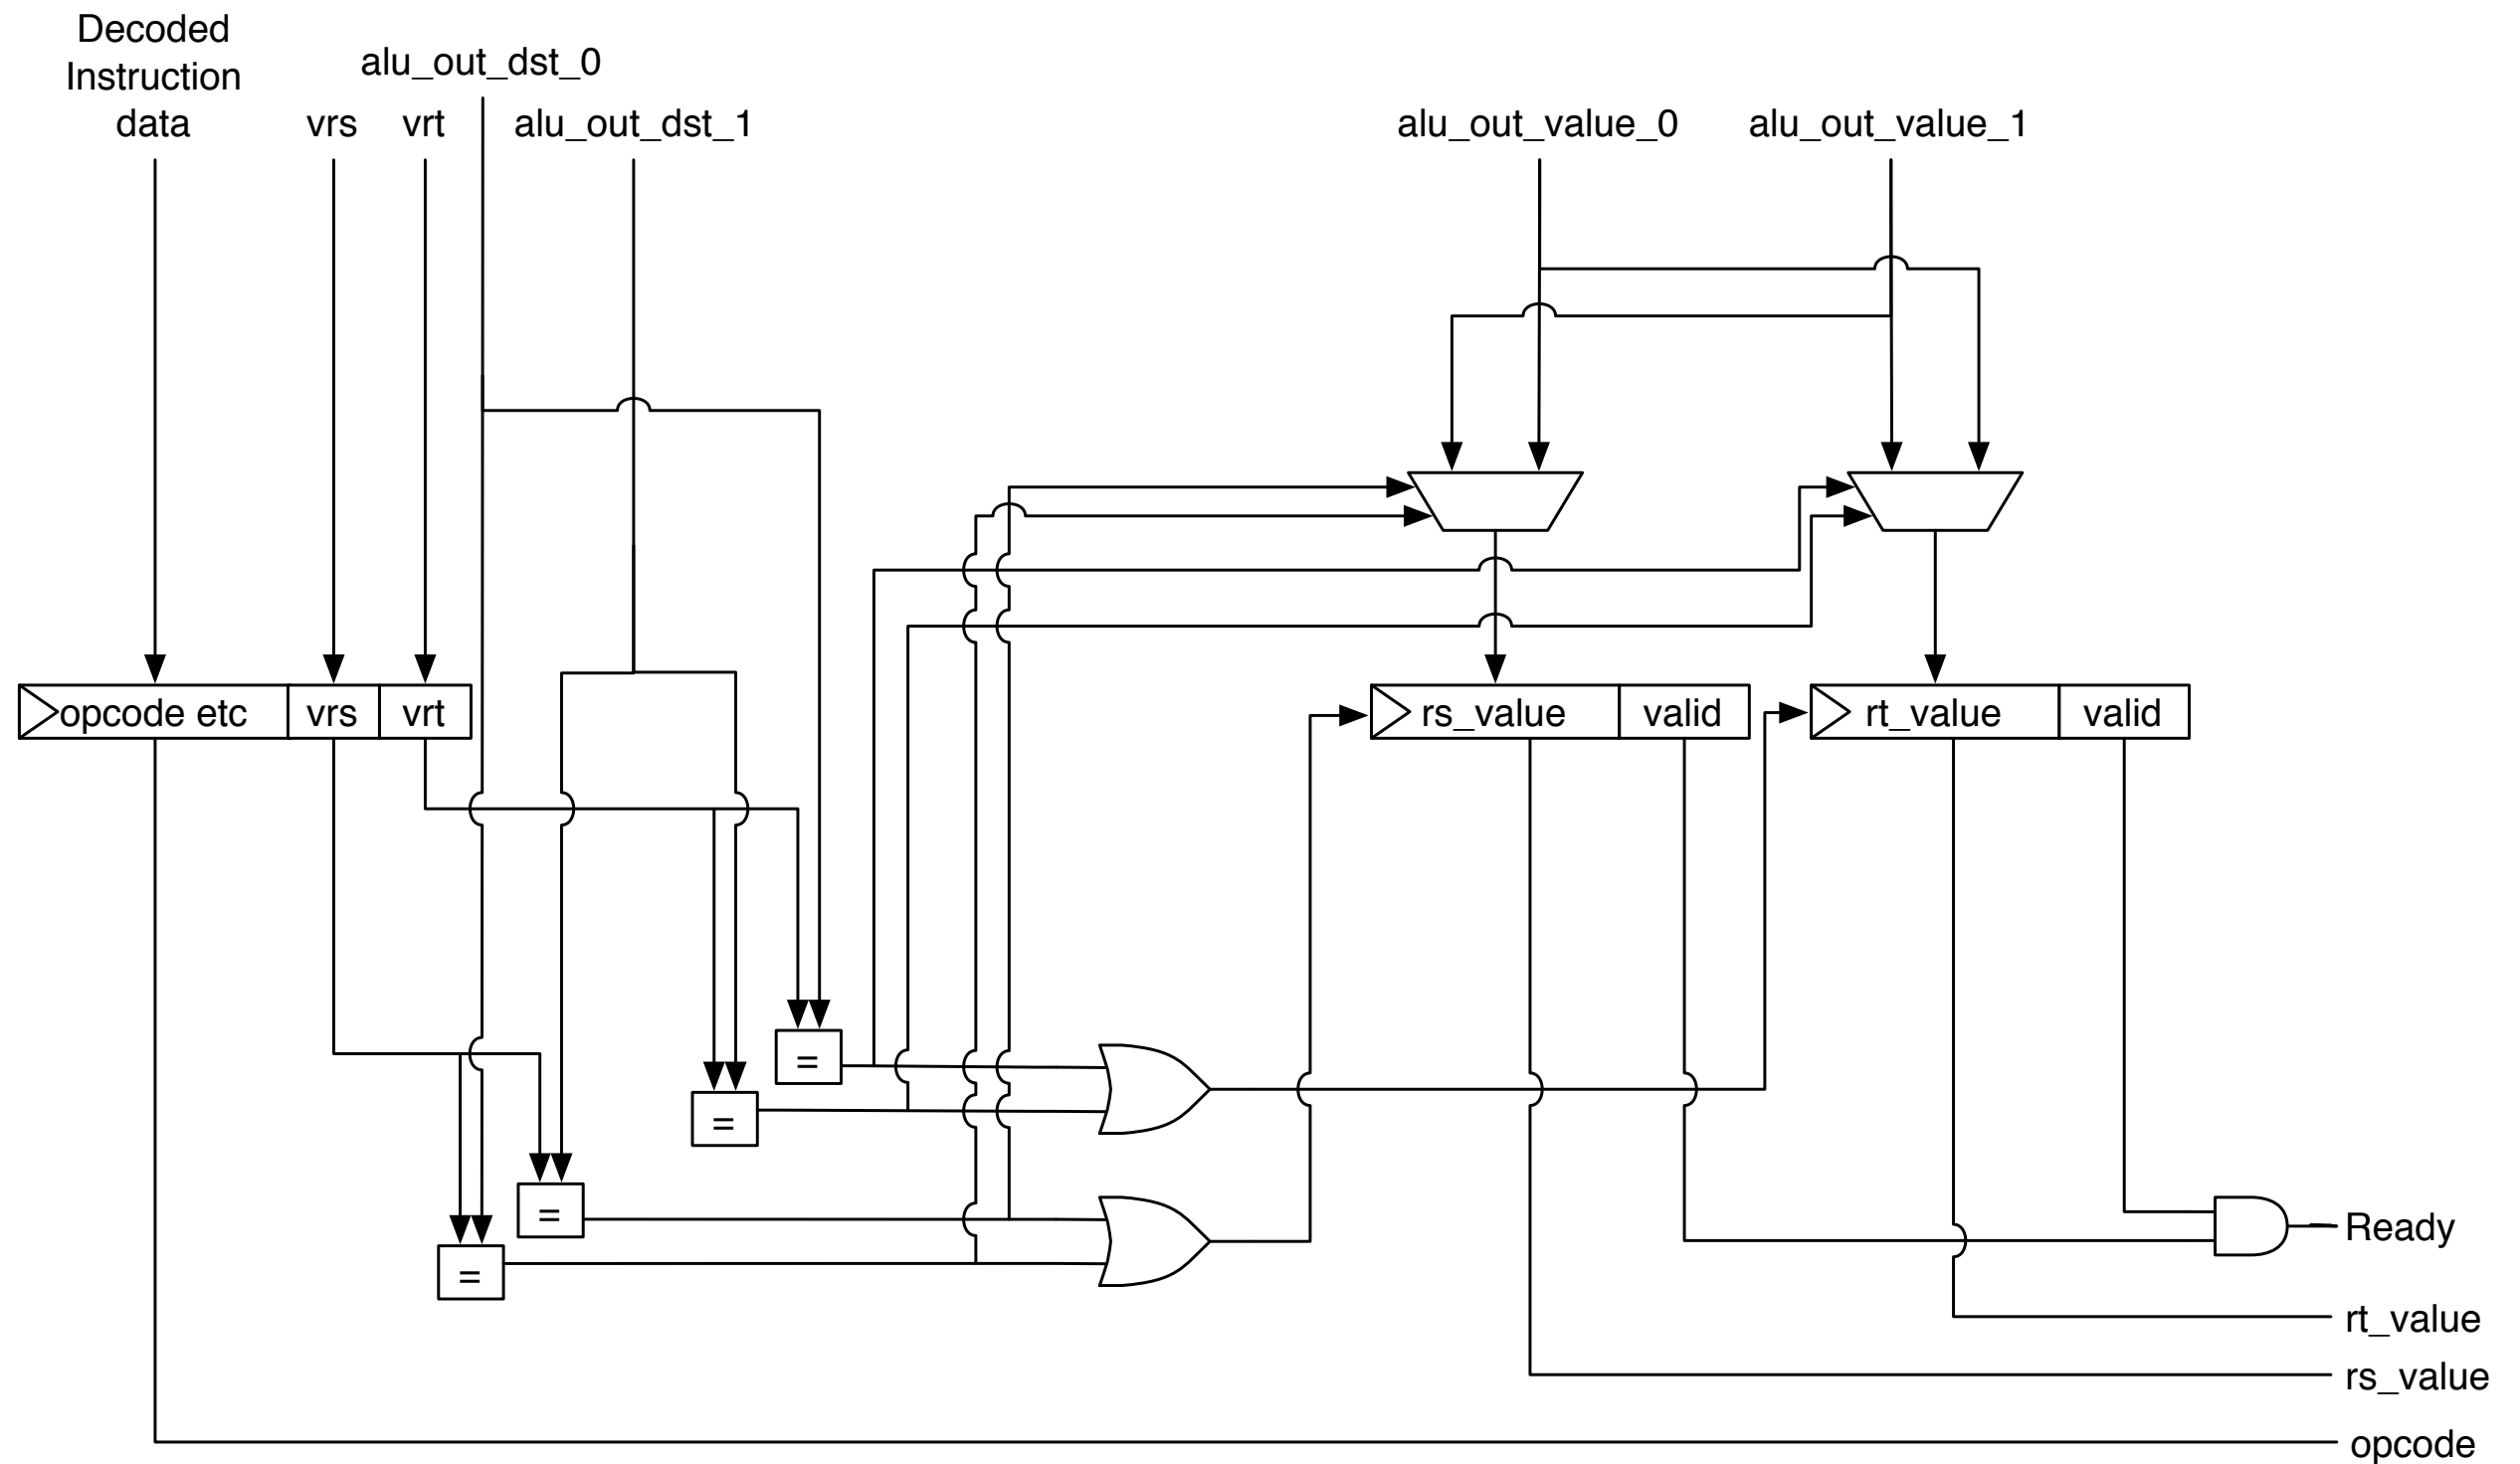
- A new “schedule” stage manages the “Instruction Window”
- The window holds the set of instruction the processor examines
  - The fetch and decode fill the window
  - Execute stage drains it
- Typically, OOO pipelines are also “wide” but it is not necessary.
- Impacts
  - More forwarding, More stalls, longer branch resolution
  - Fundamentally more work per instruction.



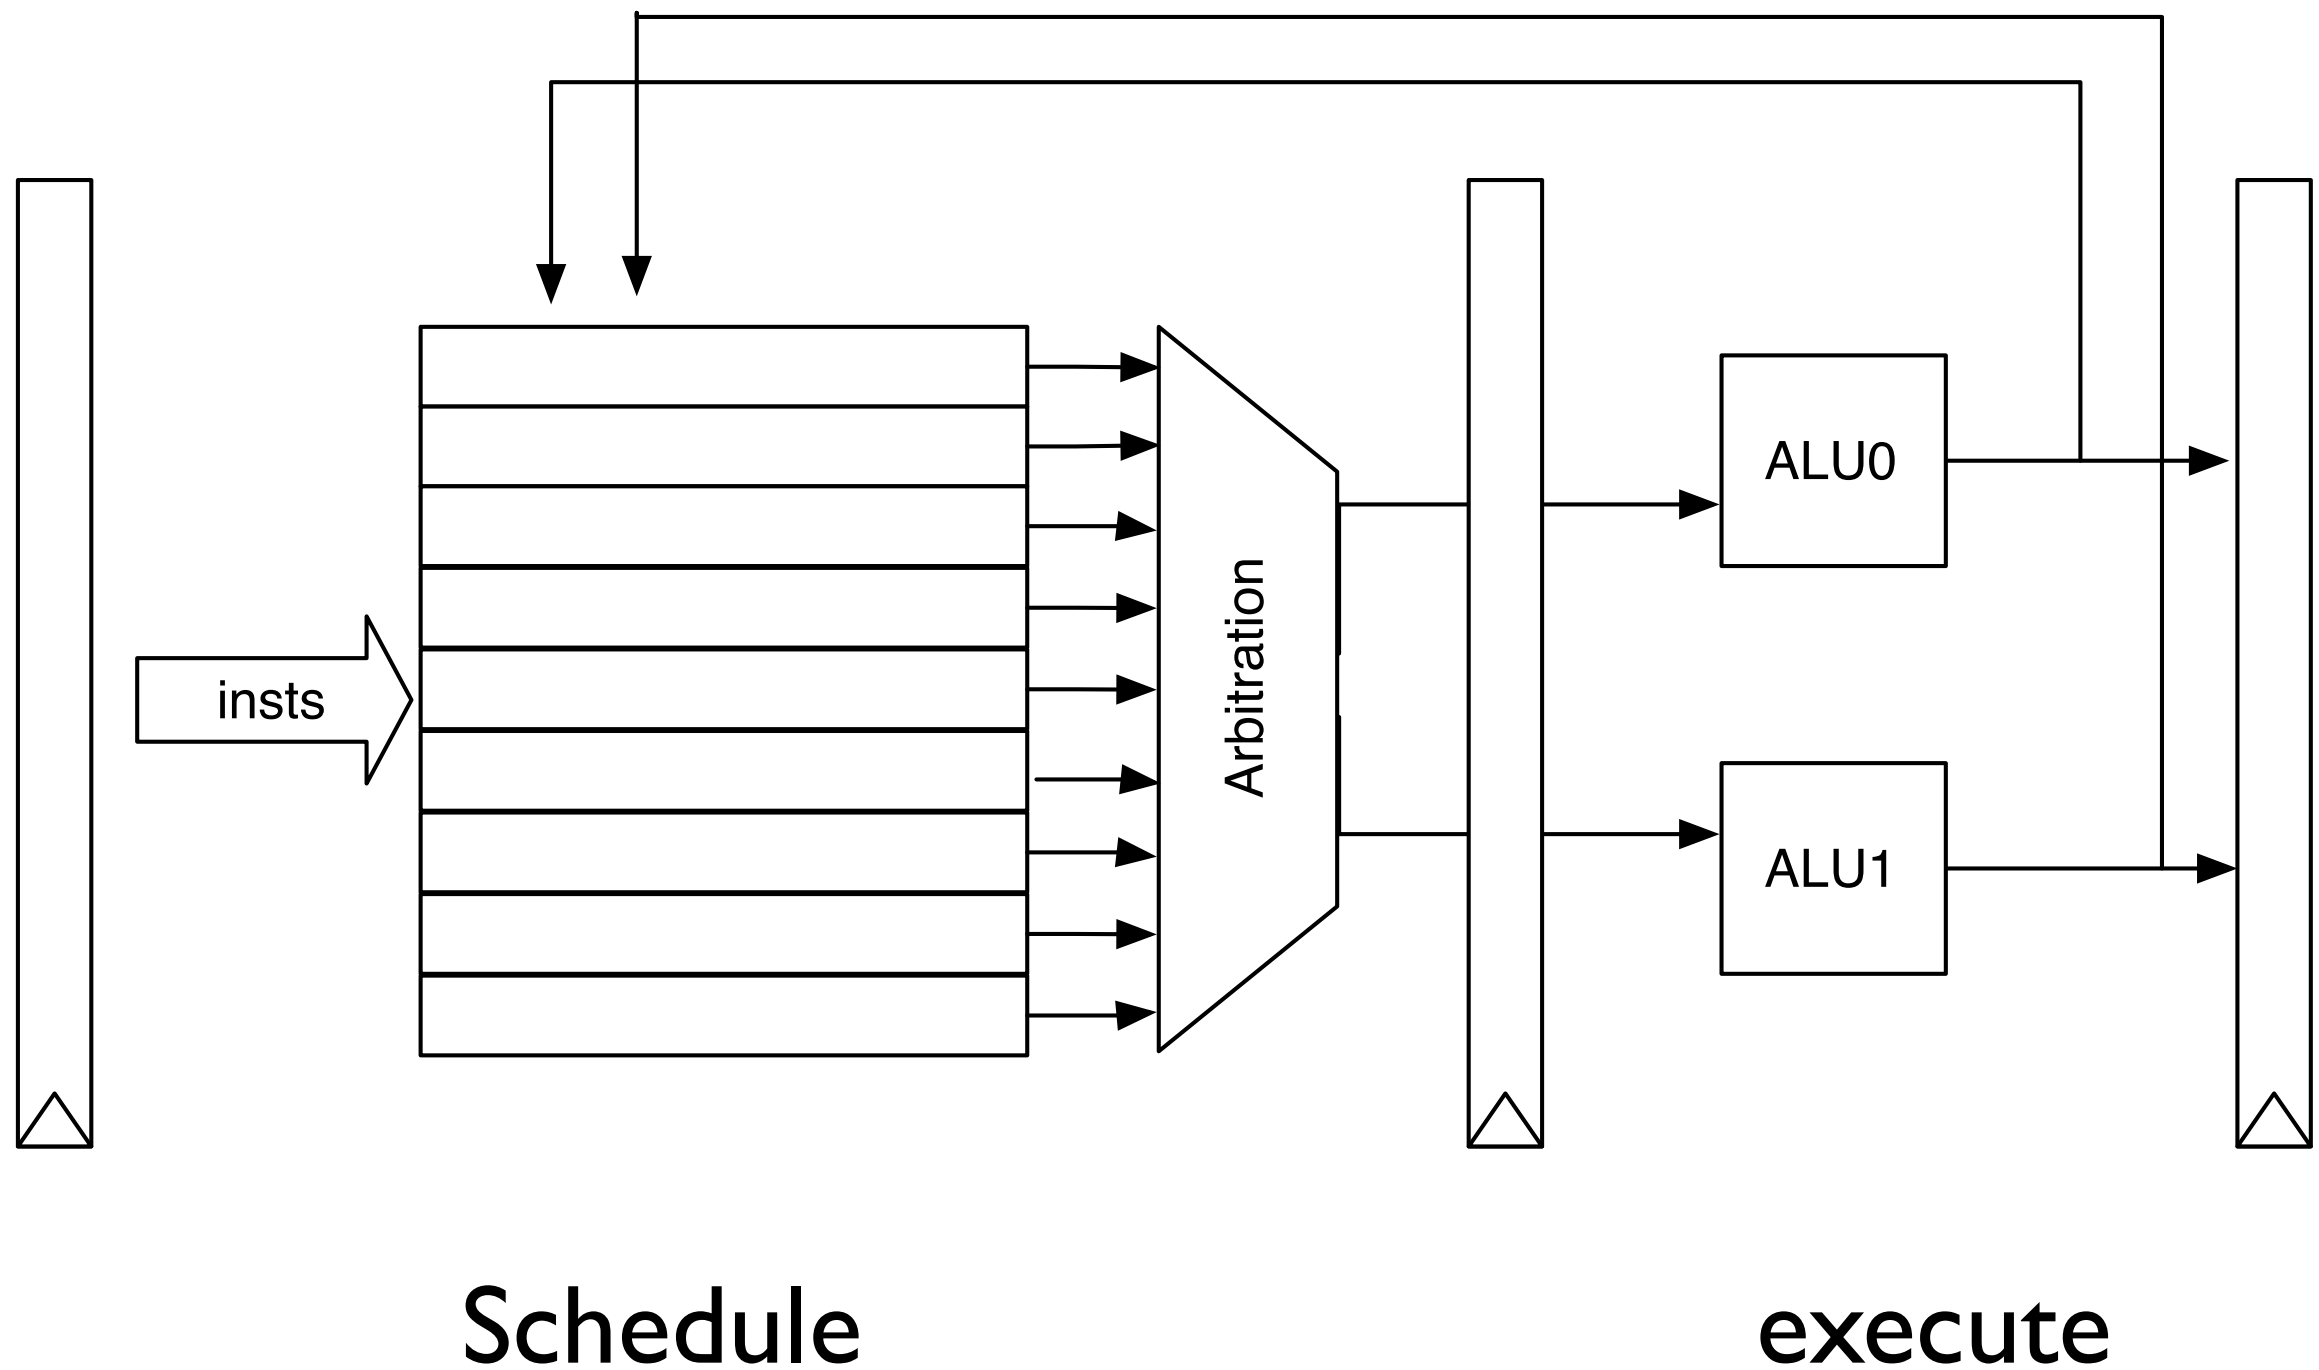
# The Instruction Window

- The “Instruction Window” is the set of instruction the processor examines
  - The fetch and decode fill the window
  - Execute stage drains it
- The larger the window, the more parallelism the processor can find, but...
- Keeping the window filled is a challenge

# The Issue Window



# The Issue Window



# Keeping the Window Filled

- Keeping the instruction window filled is key!
- Instruction windows are about 32 instructions
  - (size is limited by their complexity, which is considerable)
- Branches are every 4-5 instructions.
- This means that the processor predict 6-8 consecutive branches correctly to keep the window full.
- On a mispredict, you flush the pipeline, which includes the emptying the window.

# How Much Parallelism is There?

- Not much, in the presence of WAW and WAR dependences.
- These arise because we must reuse registers, and there are a limited number we can freely reuse.
- How can we get rid of them?

# Removing False Dependences


- If WAW and WAR dependences arise because we have too few registers
  - Let's add more!
- But! We can't! The Architecture only gives us 32 (why or why did we only use 5 bits?)
- Solution:
  - Define a set of internal "physical" register that is as large as the number of instructions that can be "in flight" -- 128 in a recent intel chip.
  - Every instruction in the pipeline gets a registers
  - Maintaining a register mapping table that determines which physical register currently holds the value for the required "architectural" registers.
- This is called "Register Renaming"

# Alpha 21264: Renaming

Register map table

```
1: Add  r3, r2, r3
2: Sub  r2, r1, r3
3: Mult r1, r3, r1
4: Add  r2, r3, r1
5: Add  r2, r1, r3
```

	r1	r2	r3
0:	p1	p2	p3
1:			
2:			
3:			
4:			
5:			






# Alpha 21264: Renaming

Register map table

```
1: Add  r3, r2, r3
2: Sub  r2, r1, r3
3: Mult r1, r3, r1
4: Add  r2, r3, r1
5: Add  r2, r1, r3
```

	r1	r2	r3
0:	p1	p2	p3
1:			
2:			
3:			
4:			
5:			

p1 currently holds the value of architectural registers r1




→ RAW → WAW → WAR


# Alpha 21264: Renaming

```
1: Add  r3, r2, r3
2: Sub  r2, r1, r3
3: Mult r1, r3, r1
4: Add  r2, r3, r1
5: Add  r2, r1, r3
```

p4, p2, p3



	r1	r2	r3
0:	p1	p2	p3
1:	p1	p2	p4
2:			
3:			
4:			
5:			




→ RAW → WAW → WAR


# Alpha 21264: Renaming

```
1: Add  r3, r2, r3
2: Sub  r2, r1, r3
3: Mult r1, r3, r1
4: Add  r2, r3, r1
5: Add  r2, r1, r3
```

```
p4, p2, p3
p5, p1, p4
```




	r1	r2	r3
0:	p1	p2	p3
1:	p1	p2	p4
2:	p1	p5	p4
3:			
4:			
5:			



→ RAW → WAW → WAR


# Alpha 21264: Renaming

1: Add r3, r2, r3  
2: Sub r2, r1, r3  
3: Mult r1, r3, r1  
4: Add r2, r3, r1  
5: Add r2, r1, r3



p4, p2, p3  
p5, p1, p4  
p6, p4, p1


	r1	r2	r3
0:	p1	p2	p3
1:	p1	p2	p4
2:	p1	p5	p4
3:	p6	p5	p4
4:			
5:			



→ RAW → WAW → WAR


# Alpha 21264: Renaming

1: Add r3, r2, r3  
2: Sub r2, r1, r3  
3: Mult r1, r3, r1  
4: Add r2, r3, r1  
5: Add r2, r1, r3



p4, p2, p3  
p5, p1, p4  
p6, p4, p1  
p7, p4, p6


	r1	r2	r3
0:	p1	p2	p3
1:	p1	p2	p4
2:	p1	p5	p4
3:	p6	p5	p4
4:	p6	p7	p4
5:			



→ RAW → WAW → WAR


# Alpha 21264: Renaming

1: Add r3, r2, r3  
2: Sub r2, r1, r3  
3: Mult r1, r3, r1  
4: Add r2, r3, r1  
5: Add r2, r1, r3



p4, p2, p3  
p5, p1, p4  
p6, p4, p1  
p7, p4, p6  
p8, p6, p4


	r1	r2	r3
0:	p1	p2	p3
1:	p1	p2	p4
2:	p1	p5	p4
3:	p6	p5	p4
4:	p6	p7	p4
5:	p6	p8	p4



→ RAW → WAW → WAR




# Alpha 21264: Renaming



1: Add r3, r2, r3  
 2: Sub r2, r1, r3  
 3: Mult r1, r3, r1  
 4: Add r2, r3, r1  
 5: Add r2, r1, r3



p4, p2, p3  
 p5, p1, p4  
 p6, p4, p1  
 p7, p4, p6  
 p8, p6, p4

	r1	r2	r3
0:	p1	p2	p3
1:	p1	p2	p4
2:	p1	p5	p4
3:	p6	p5	p4
4:	p6	p7	p4
5:	p6	p8	p4



 RAW
  WAW
  WAR

### Question 1: Multiple Answer

Average Score 1.48538 points

Which of the following can significantly reduce conflict misses?

Correct	Answers	Percent Correct	Percent Incorrect
	Reducing associativity.	96.491%	3.509%
✓	Increasing associativity.	94.737%	5.263%
✓	Increasing the number of cache lines.	80.702%	19.298%
✓	Improving the replacement policy.	50.877%	49.123%



## Question 2: Multiple Answer

Average Score 0.5614 points

Which of the following can significantly reduced capacity misses?

Correct Answers

Percent Correct Percent Incorrect

Increasing associativity.

40.351%

59.649%



Increasing the number of cache lines, while holding the cache line size constant.

80.702%

19.298%

Improving the cache replacement policy.

85.965%

14.035%

Making the cache write through, rather than write back.


87.719%

12.281%

### Question 3: True/False

Average Score 0.96491 points

Write back is always better than write through.

Correct	Answers	Percent Answered
	True	3.509%
	False	96.491%
	<i>Unanswered</i>	0%

### Question 5: Calculated Numeric

Average Score 1.21053 points

Compute the average CPI for a system with an L1 hit rate of 90% and an L2 hit rate of 95%. Assume that L1 hits take 1 cycle, L2 hits take 20 cycles, and main memory accesses take 100 cycles. Further assume that 20% of instructions are memory accesses, and that the processor has a perfect branch predictor.

- $CPI = 0.8 * 1 + //$  non-memory
- $0.2 * //$  memory
  - $(.9 * 1 //$  L1 hits
  - $+ 0.1 * //$  L1 misses
    - $(0.95 * 20 //$  L2 hits
    - $+ 0.05 * 100) //$  L2 misses
- If you answered 1.66 or 1.8 let me know.

### Question 6: True/False

Average Score 0.57895 points

Increasing the number of sets in an associative cache causes the number of index bits to increase.

Correct	Answers	Percent Answered
<input checked="" type="checkbox"/>	True	57.895%
	False	42.105%
	<i>Unanswered</i>	0%

### Question 7: Multiple Choice

Average Score 2.89474 points

Given a 32 KB, 2-way set associative, 64 byte cache lines, give the number of offset, index, and tag bits required to access it.

Correct

Percent Answered

Index:32  
offset:8  
Tag:12

0%

Index:7  
offset: 5  
Tag:19

0%

Index:9  
offset:5  
Tag:17

3.509%



Index:8  
offset:6  
Tag:18


96.491%

*Unanswered*

0%

# New OOO Pipeline

- The register file is larger (to hold the physical registers)
- The pipeline is longer
  - more forwarding
  - Longer branch delay
- The payoff had better be significant (and it is)



# Modern OOO Processors

- The fastest machines in the world are OOO superscalars
- AMD Barcelona
  - 6-wide issue
  - 106 instructions in flight at once.
- Intel Nehalem
  - 5-way issue to 12 ALUs
  - > 128 instructions in flight
- OOO provides the most benefit for memory operations.
  - Non-dependent instructions can keep executing during cache misses.
  - This is so-called “memory-level parallelism.”
  - It is enormously important. CPU performance is (almost) all about memory performance nowadays (remember the memory wall graphs!)

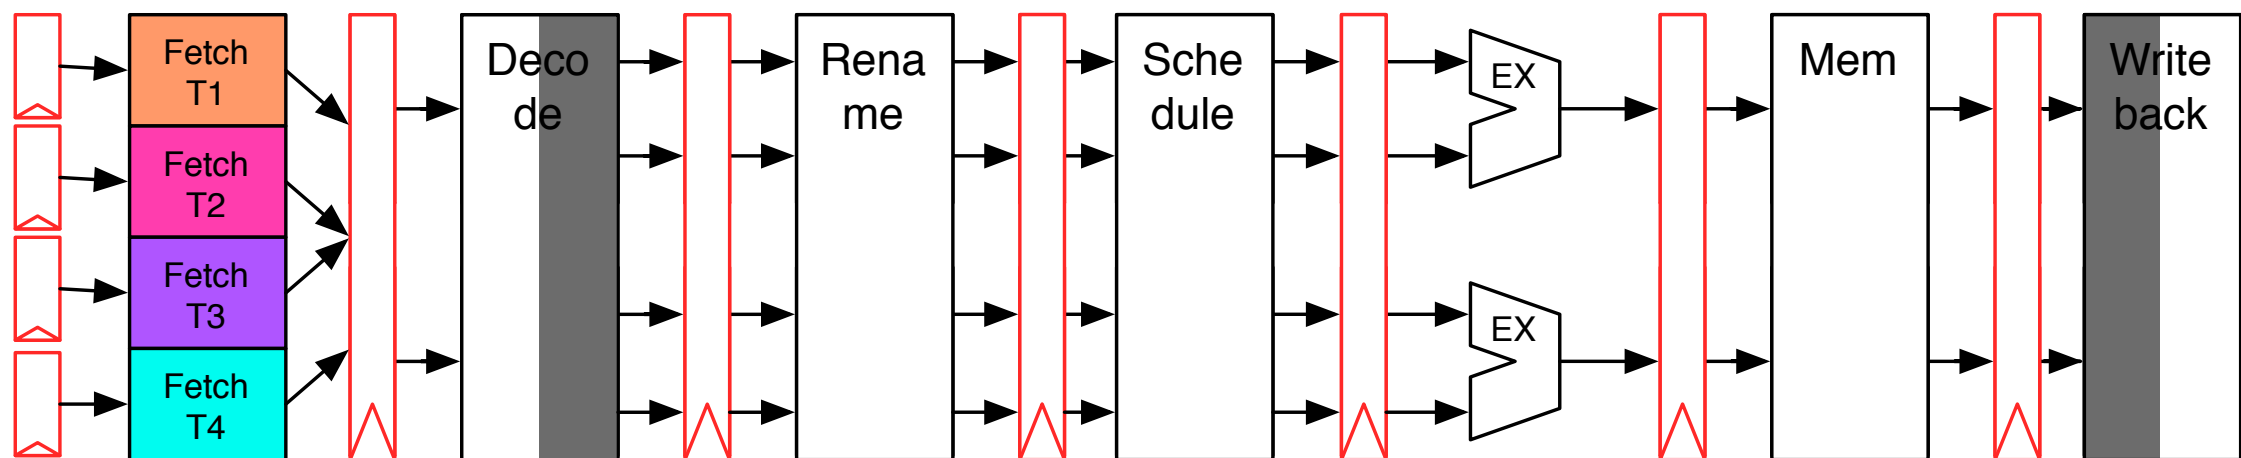
# The Problem with OOO

- Even the fastest OOO machines only get about 1-2 IPC, even though they are 4-5 wide.
- Problems
  - Insufficient ILP within applications. -- 1-2 per thread, usually
  - Poor branch prediction performance
  - Single threads also have little memory parallelism.
- Observation
  - On many cycles, many ALUs and instruction queue slots sit empty



# Simultaneous Multithreading

- AKA HyperThreading in Intel machines
- Run multiple threads at the same time
- Just throw all the instructions into the pipeline
- Keep some separate data for each
  - Renaming table
  - TLB entries
  - PCs
- But the rest of the hardware is shared.
- It is surprisingly simple (but still quite complicated)



# SMT Advantages

- Exploit the ILP of multiple threads at once
  - Less dependence on branch prediction (fewer correct predictions required per thread)
  - Less idle hardware (increased power efficiency)
  - Much higher IPC -- up to 4
  - Disadvantages: threads can fight over resources and slow each other down.
- 
- Historical footnote: Invented, in part, by our own Dean Tullsen when he was at UW