

DRAM

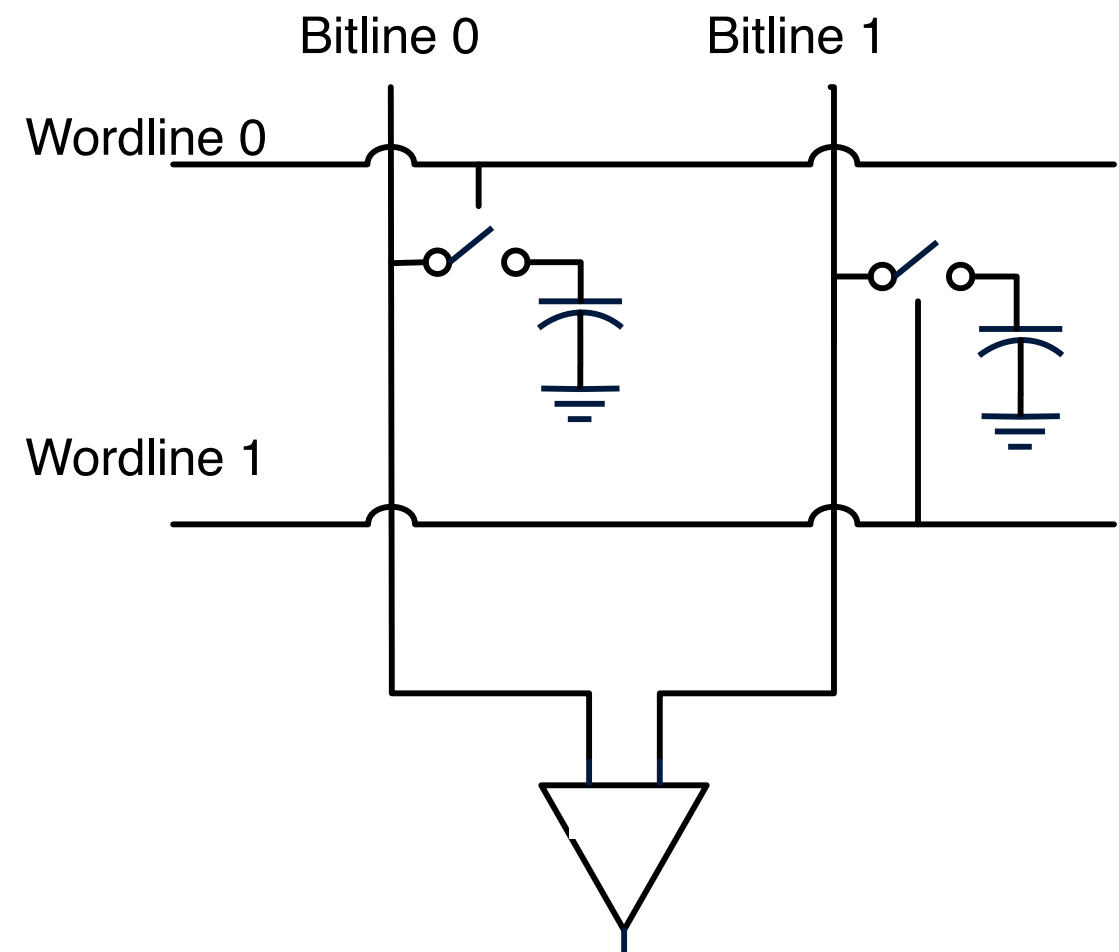
Dynamic Random Access Memory (DRAM)

- Storage
 - Charge on a capacitor
 - Decays over time (us-scale)
 - This is the “dynamic” part.
 - About $6F^2$: 20x better than SRAM
- Reading
 - Precharge
 - Assert word line
 - Sense output
 - Refresh data

Only one bit line is read at a time.
The other bit line serves as a reference.
The bit cells attached to Wordline 1 are not shown.

Dynamic Random Access Memory (DRAM)

- Storage
 - Charge on a capacitor
 - Decays over time (us-scale)
 - This is the “dynamic” part.
 - About $6F^2$: 20x better than SRAM
- Reading
 - Precharge
 - Assert word line
 - Sense output
 - Refresh data



Only one bit line is read at a time.
The other bit line serves as a reference.
The bit cells attached to Wordline 1 are not shown.

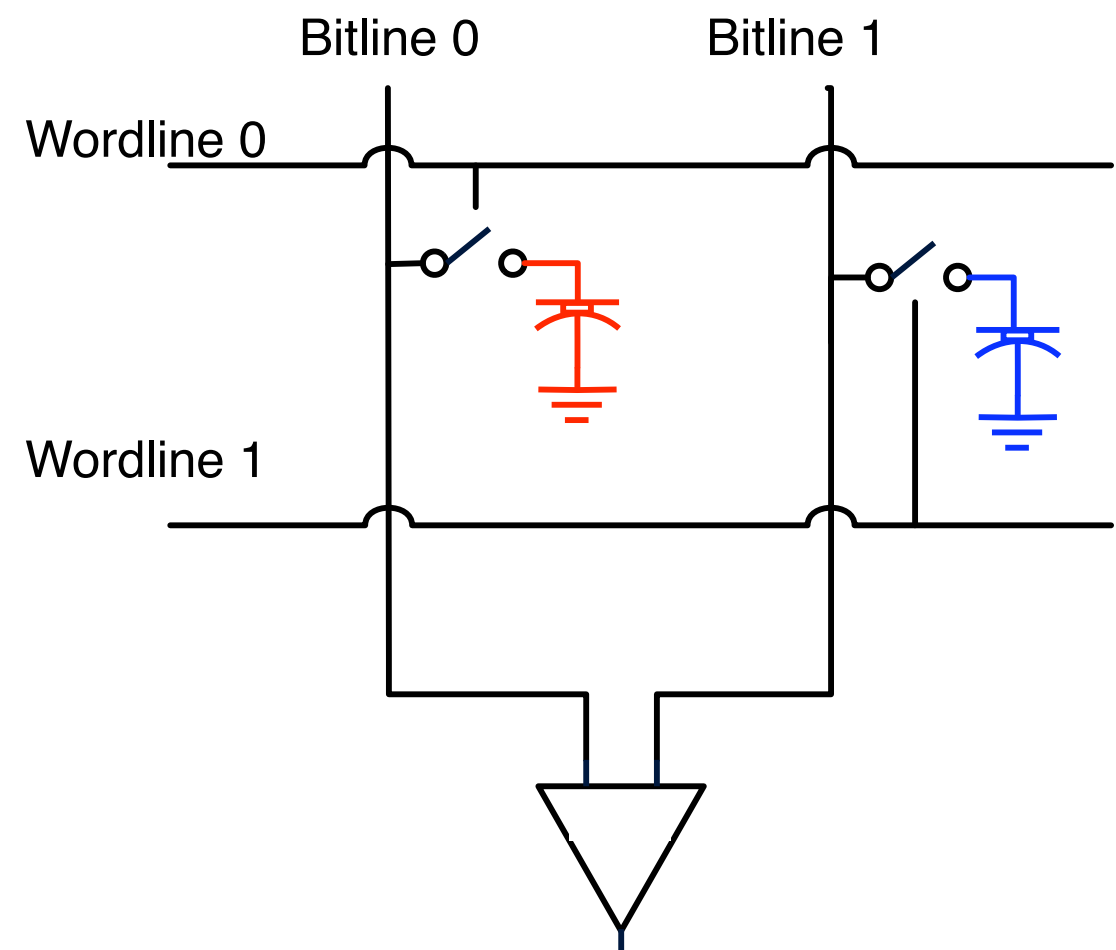
Dynamic Random Access Memory (DRAM)

- Storage

- Charge on a capacitor
- Decays over time (us-scale)
- This is the “dynamic” part.
- About $6F^2$: 20x better than SRAM

- Reading

- Precharge
- Assert word line
- Sense output
- Refresh data



Only one bit line is read at a time.
The other bit line serves as a reference.
The bit cells attached to Wordline 1 are not shown.

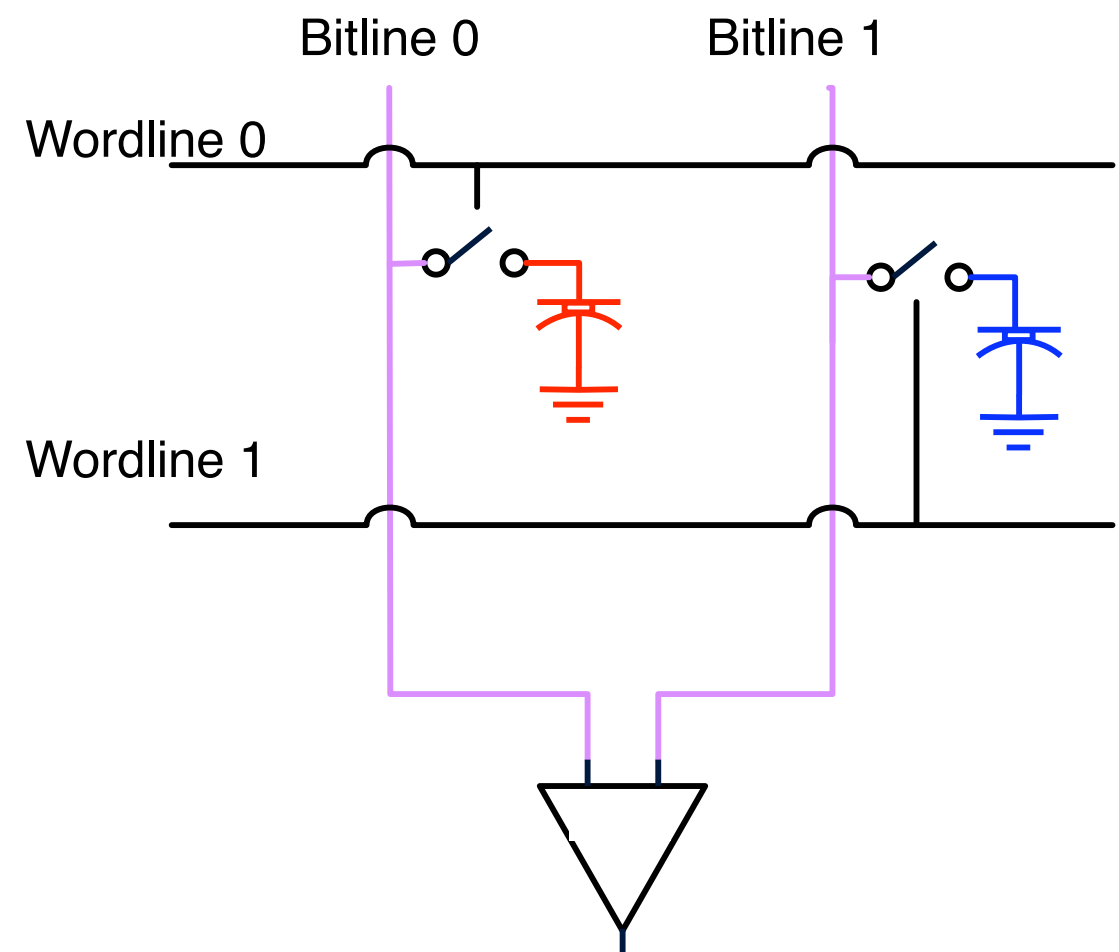
Dynamic Random Access Memory (DRAM)

- Storage

- Charge on a capacitor
- Decays over time (us-scale)
- This is the “dynamic” part.
- About $6F^2$: 20x better than SRAM

- Reading

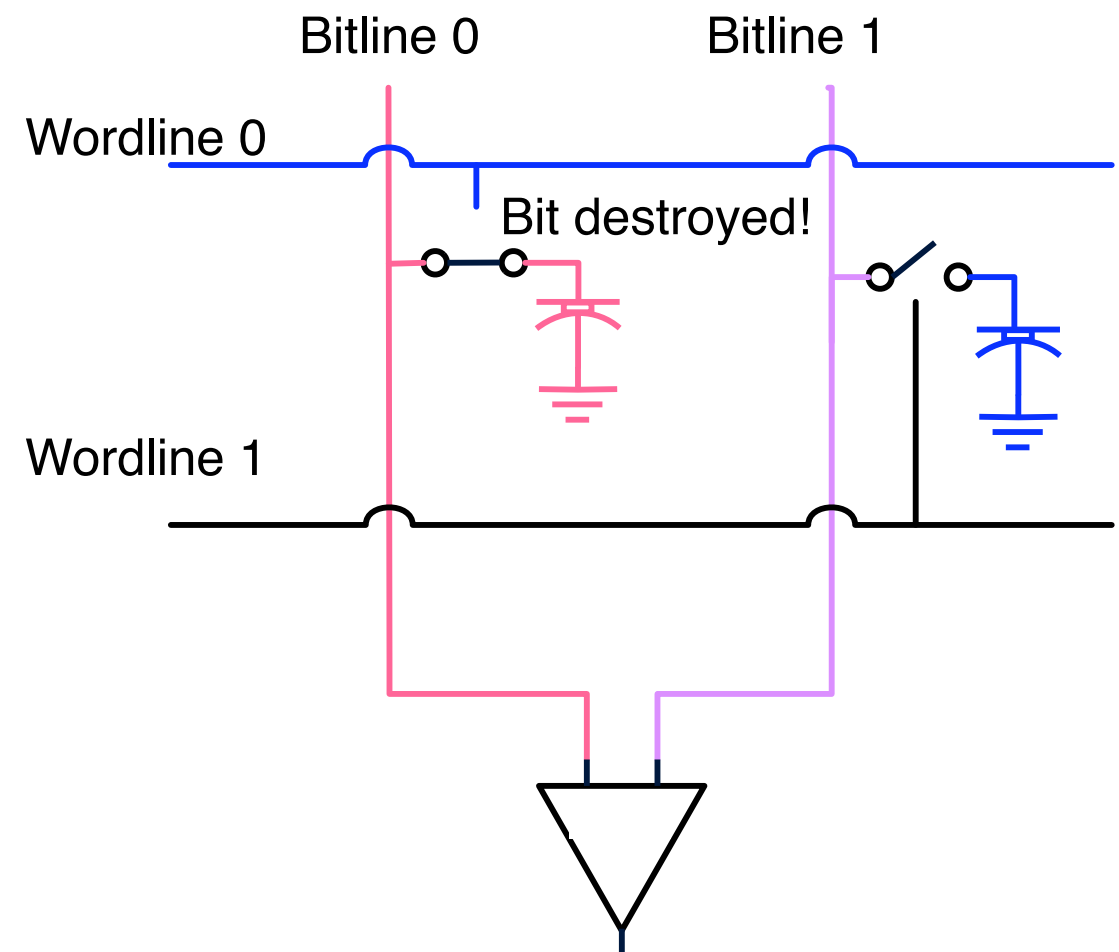
- Precharge
- Assert word line
- Sense output
- Refresh data



Only one bit line is read at a time.
The other bit line serves as a reference.
The bit cells attached to Wordline 1 are not shown.

Dynamic Random Access Memory (DRAM)

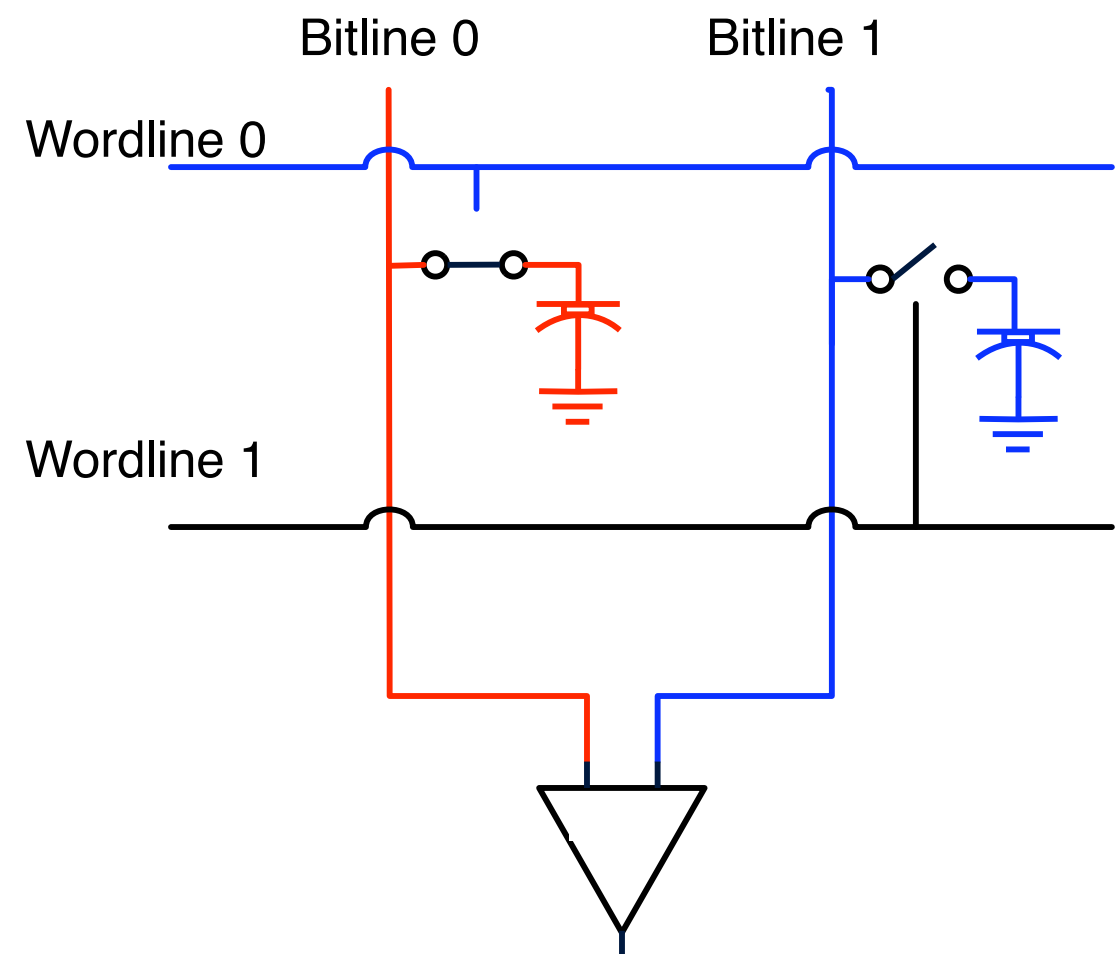
- Storage
 - Charge on a capacitor
 - Decays over time (us-scale)
 - This is the “dynamic” part.
 - About $6F^2$: 20x better than SRAM
- Reading
 - Precharge
 - Assert word line
 - Sense output
 - Refresh data



Only one bit line is read at a time.
The other bit line serves as a reference.
The bit cells attached to Wordline 1 are not shown.

Dynamic Random Access Memory (DRAM)

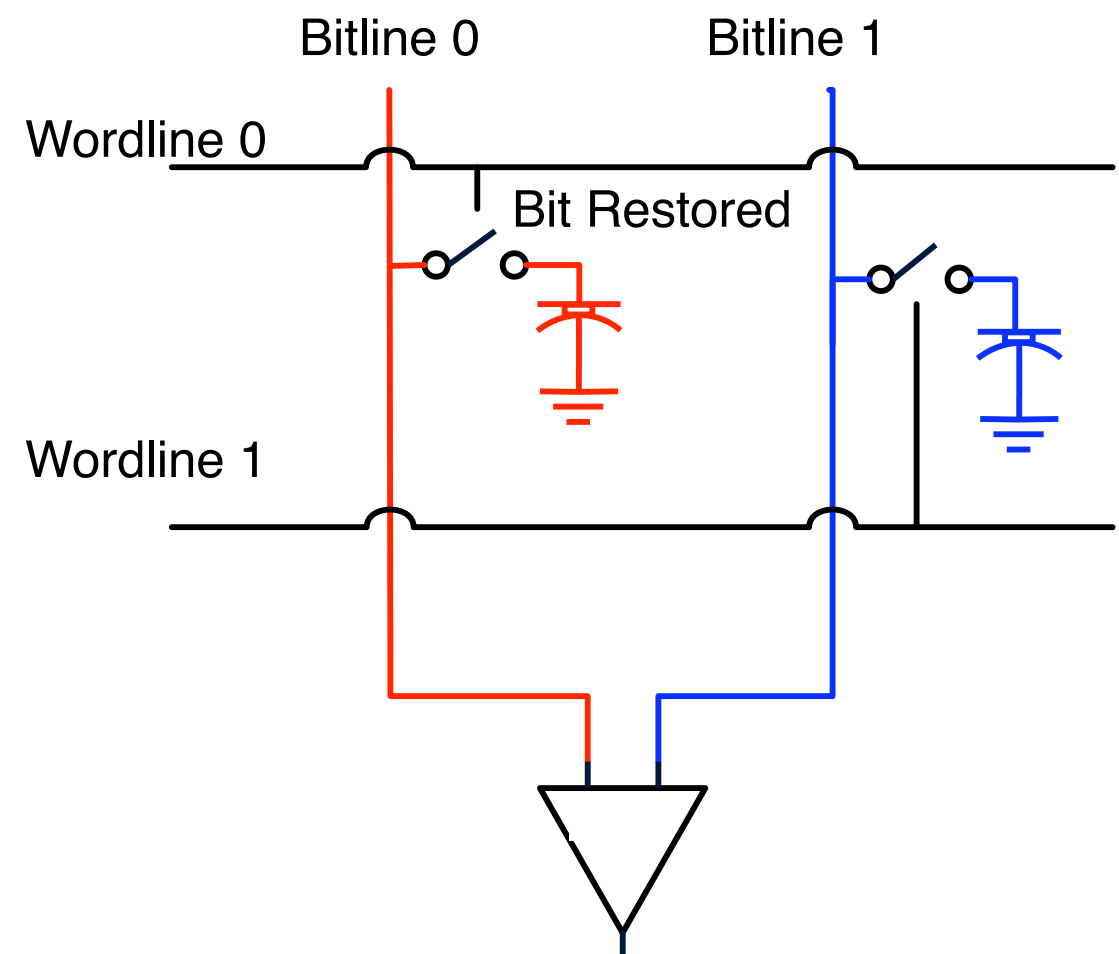
- Storage
 - Charge on a capacitor
 - Decays over time (us-scale)
 - This is the “dynamic” part.
 - About $6F^2$: 20x better than SRAM
- Reading
 - Precharge
 - Assert word line
 - Sense output
 - Refresh data



Only one bit line is read at a time.
The other bit line serves as a reference.
The bit cells attached to Wordline 1 are not shown.

Dynamic Random Access Memory (DRAM)

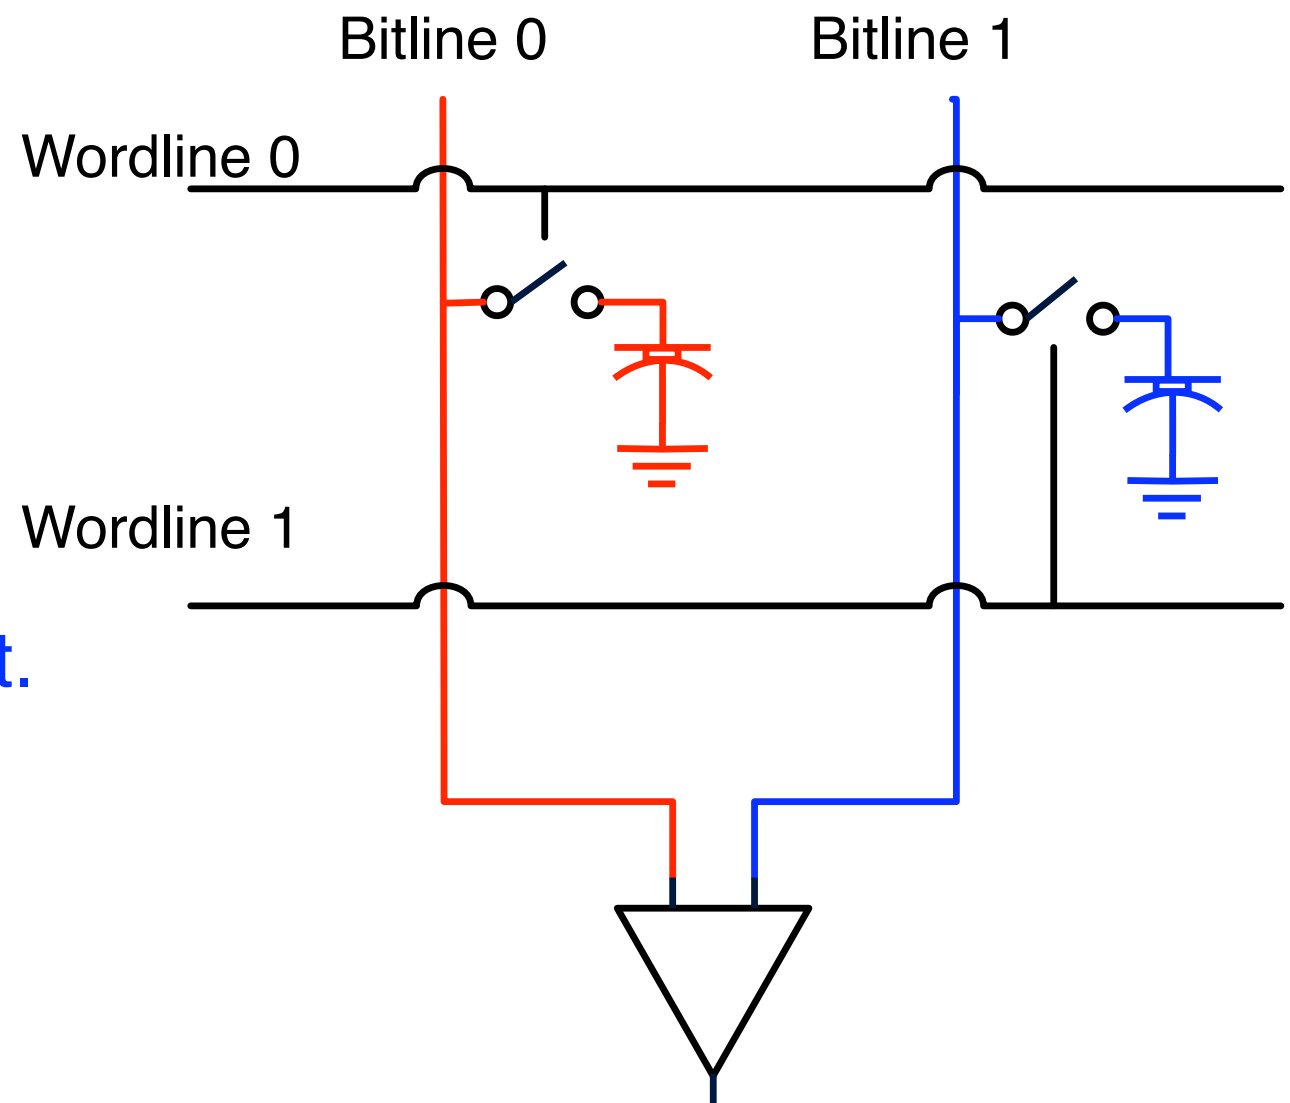
- Storage
 - Charge on a capacitor
 - Decays over time (us-scale)
 - This is the “dynamic” part.
 - About $6F^2$: 20x better than SRAM
- Reading
 - Precharge
 - Assert word line
 - Sense output
 - Refresh data




Only one bit line is read at a time.
The other bit line serves as a reference.
The bit cells attached to Wordline 1 are not shown.

DRAM: Write and Refresh

- Writing
 - Turn on the wordline
 - Override the sense amp.
- Refresh
 - Every few milli-seconds, read and re-write every bit.
 - Consumes power
 - Takes time




DRAM Lithography




Source: Hitachi/VCE, "Memory 1997"

20766




Accessing DRAM

- Apply the row address
 - “opens a page”
 - Slow (~12ns read + 24 ns precharge)
 - Contents in a “row buffer”
- Apply one or more column addrs
 - fast (~3ns)
 - Reads and/or writes




Accessing DRAM

- Apply the row address
 - “opens a page”
 - Slow (~12ns read + 24 ns precharge)
 - Contents in a “row buffer”
- Apply one or more column addrs
 - fast (~3ns)
 - Reads and/or writes




Accessing DRAM

- Apply the row address
 - “opens a page”
 - Slow (~12ns read + 24 ns precharge)
 - Contents in a “row buffer”
- Apply one or more column addrs
 - fast (~3ns)
 - Reads and/or writes




Accessing DRAM

- Apply the row address
 - “opens a page”
 - Slow (~12ns read + 24 ns precharge)
 - Contents in a “row buffer”
- Apply one or more column addrs
 - fast (~3ns)
 - Reads and/or writes



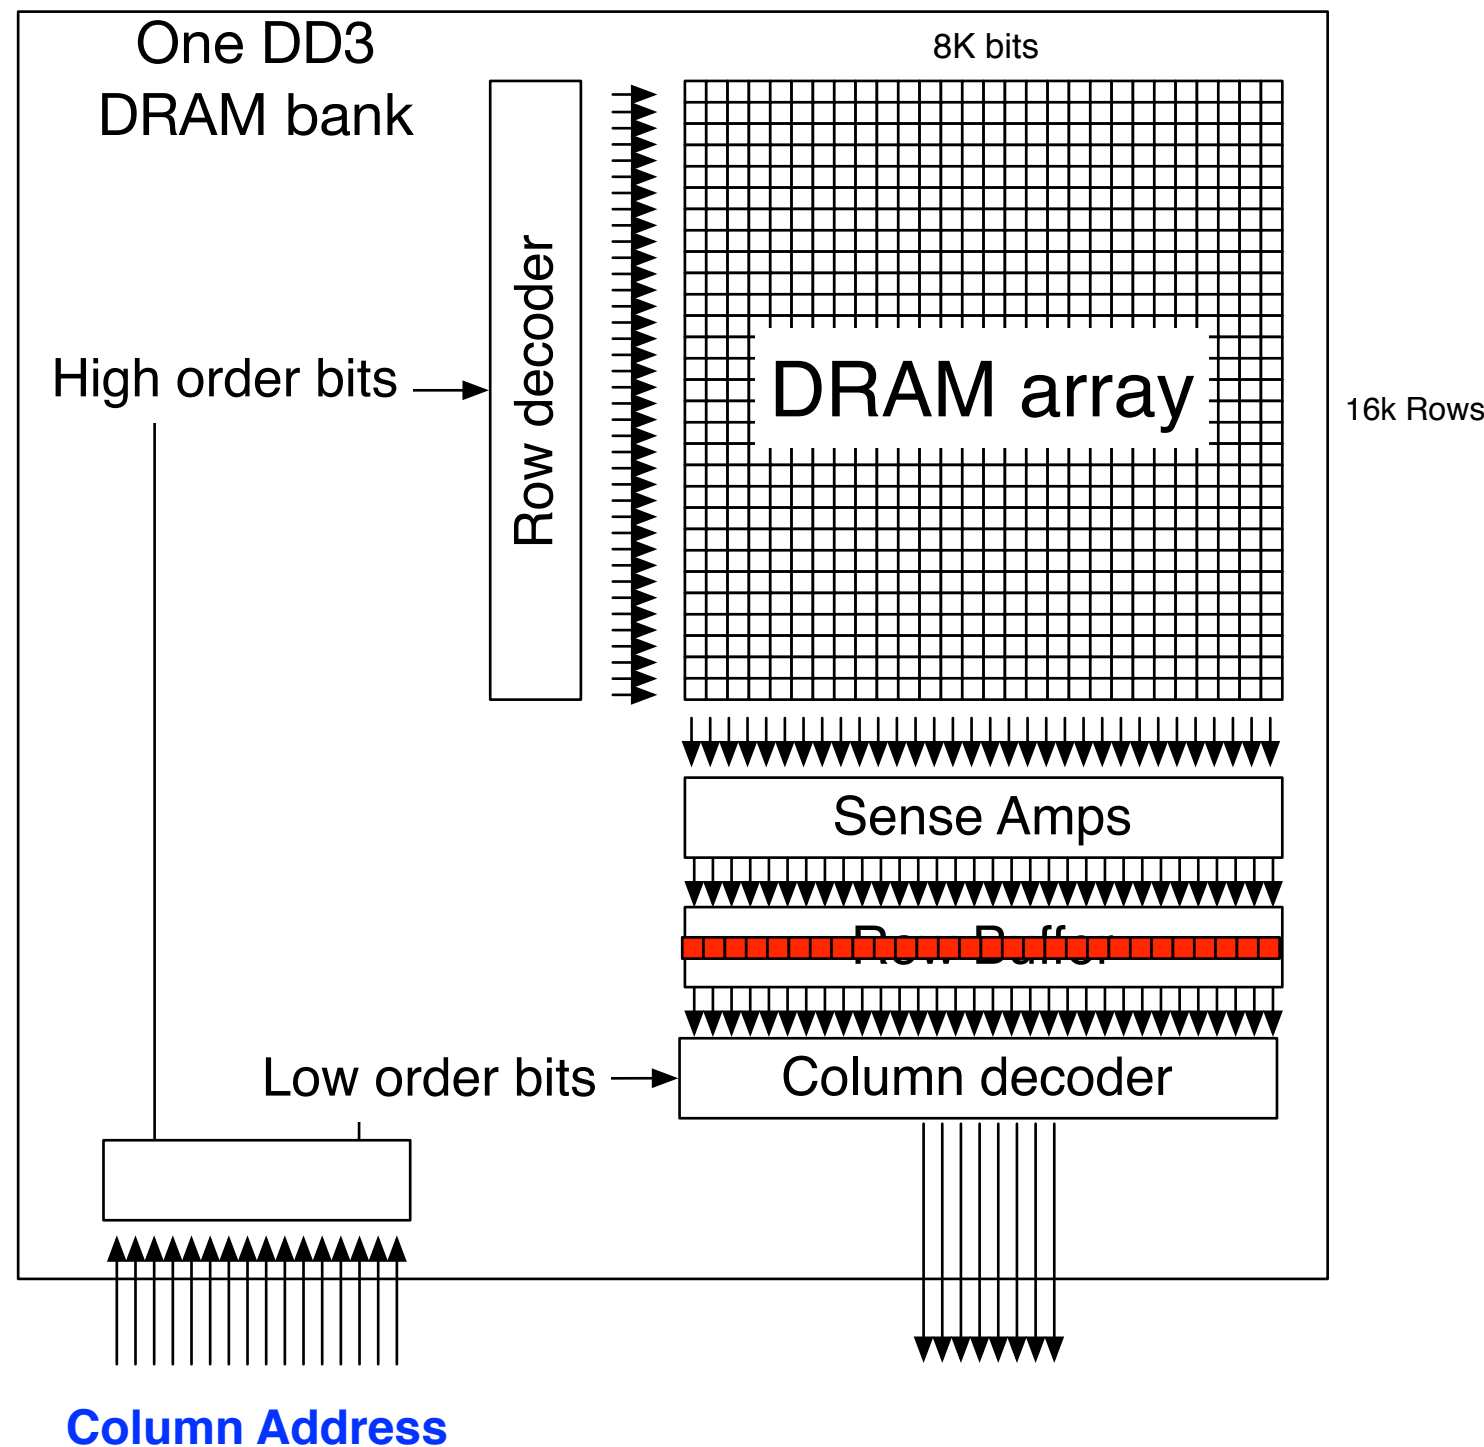
Accessing DRAM

- Apply the row address
 - “opens a page”
 - Slow (~12ns read + 24 ns precharge)
 - Contents in a “row buffer”
- Apply one or more column addrs
 - fast (~3ns)
 - Reads and/or writes



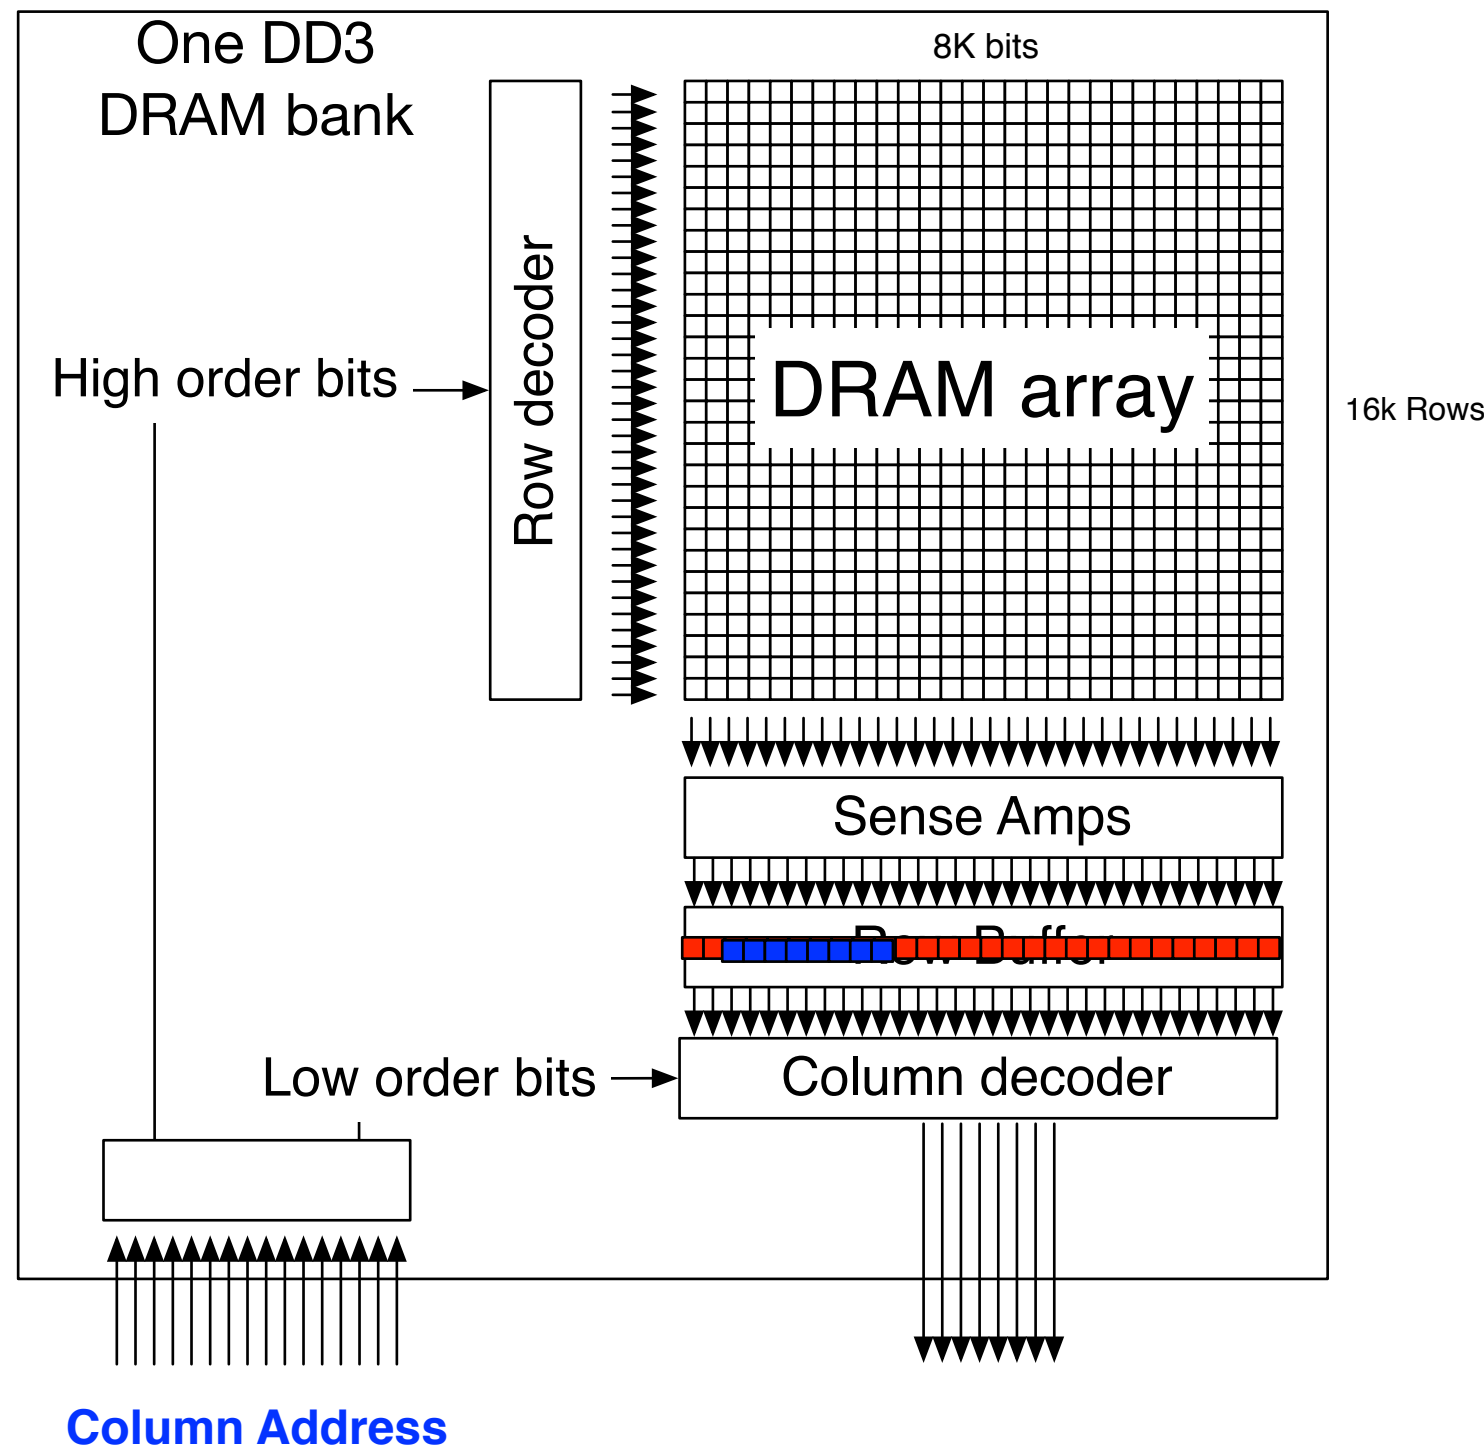
Accessing DRAM

- Apply the row address
 - “opens a page”
 - Slow (~12ns read + 24 ns precharge)
 - Contents in a “row buffer”
- Apply one or more column addrs
 - fast (~3ns)
 - Reads and/or writes



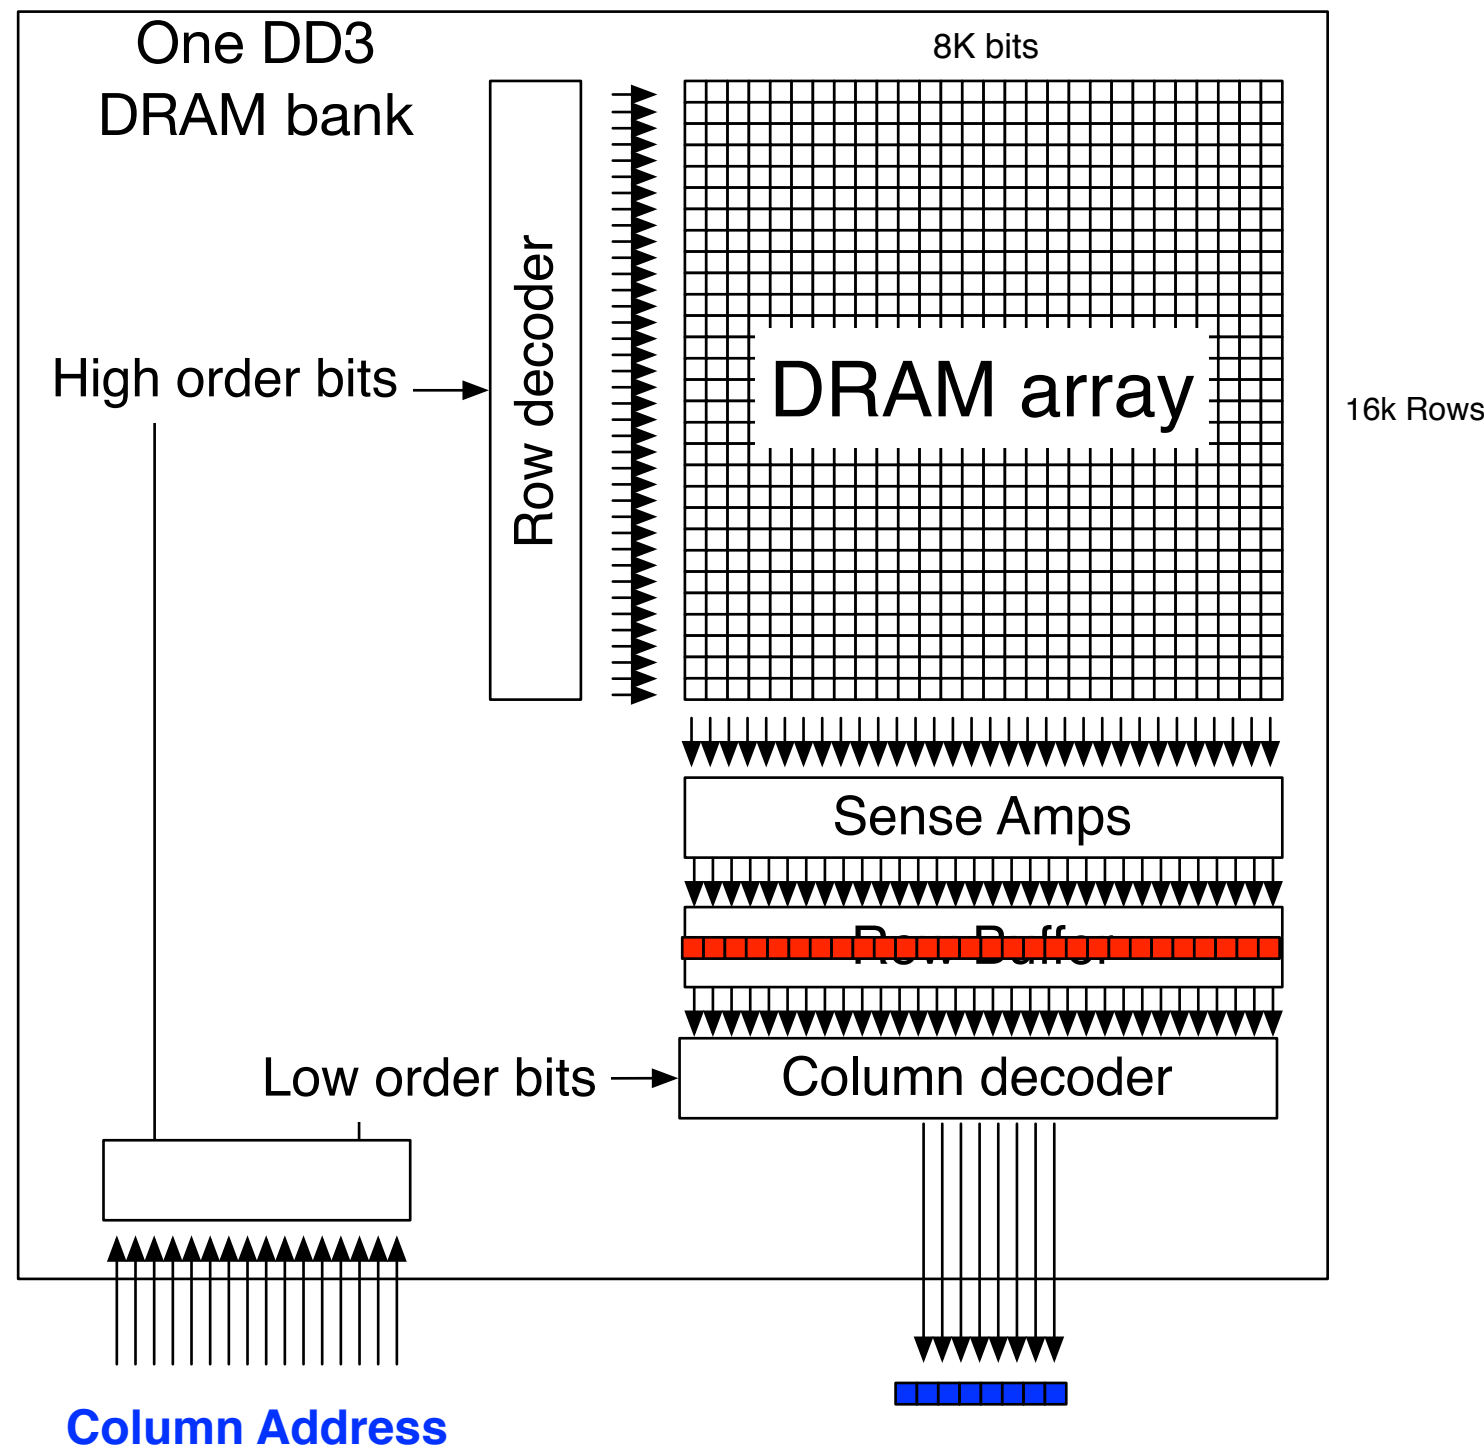
Accessing DRAM

- Apply the row address
 - “opens a page”
 - Slow (~12ns read + 24 ns precharge)
 - Contents in a “row buffer”
- Apply one or more column addrs
 - fast (~3ns)
 - Reads and/or writes




Accessing DRAM

- Apply the row address
 - “opens a page”
 - Slow (~12ns read + 24 ns precharge)
 - Contents in a “row buffer”
- Apply one or more column addrs
 - fast (~3ns)
 - Reads and/or writes




Accessing DRAM

- Apply the row address
 - “opens a page”
 - Slow (~12ns read + 24 ns precharge)
 - Contents in a “row buffer”
- Apply one or more column addrs
 - fast (~3ns)
 - Reads and/or writes




Accessing DRAM

- Apply the row address
 - “opens a page”
 - Slow (~12ns read + 24 ns precharge)
 - Contents in a “row buffer”
- Apply one or more column addrs
 - fast (~3ns)
 - Reads and/or writes




Accessing DRAM

- Apply the row address
 - “opens a page”
 - Slow (~12ns read + 24 ns precharge)
 - Contents in a “row buffer”
- Apply one or more column addrs
 - fast (~3ns)
 - Reads and/or writes




Accessing DRAM

- Apply the row address
 - “opens a page”
 - Slow (~12ns read + 24 ns precharge)
 - Contents in a “row buffer”
- Apply one or more column addrs
 - fast (~3ns)
 - Reads and/or writes




Accessing DRAM

- Apply the row address
 - “opens a page”
 - Slow (~12ns read + 24 ns precharge)
 - Contents in a “row buffer”
- Apply one or more column addrs
 - fast (~3ns)
 - Reads and/or writes




Accessing DRAM

- Apply the row address
 - “opens a page”
 - Slow (~12ns read + 24 ns precharge)
 - Contents in a “row buffer”
- Apply one or more column addrs
 - fast (~3ns)
 - Reads and/or writes




Accessing DRAM

- Apply the row address
 - “opens a page”
 - Slow (~12ns read + 24 ns precharge)
 - Contents in a “row buffer”
- Apply one or more column addrs
 - fast (~3ns)
 - Reads and/or writes



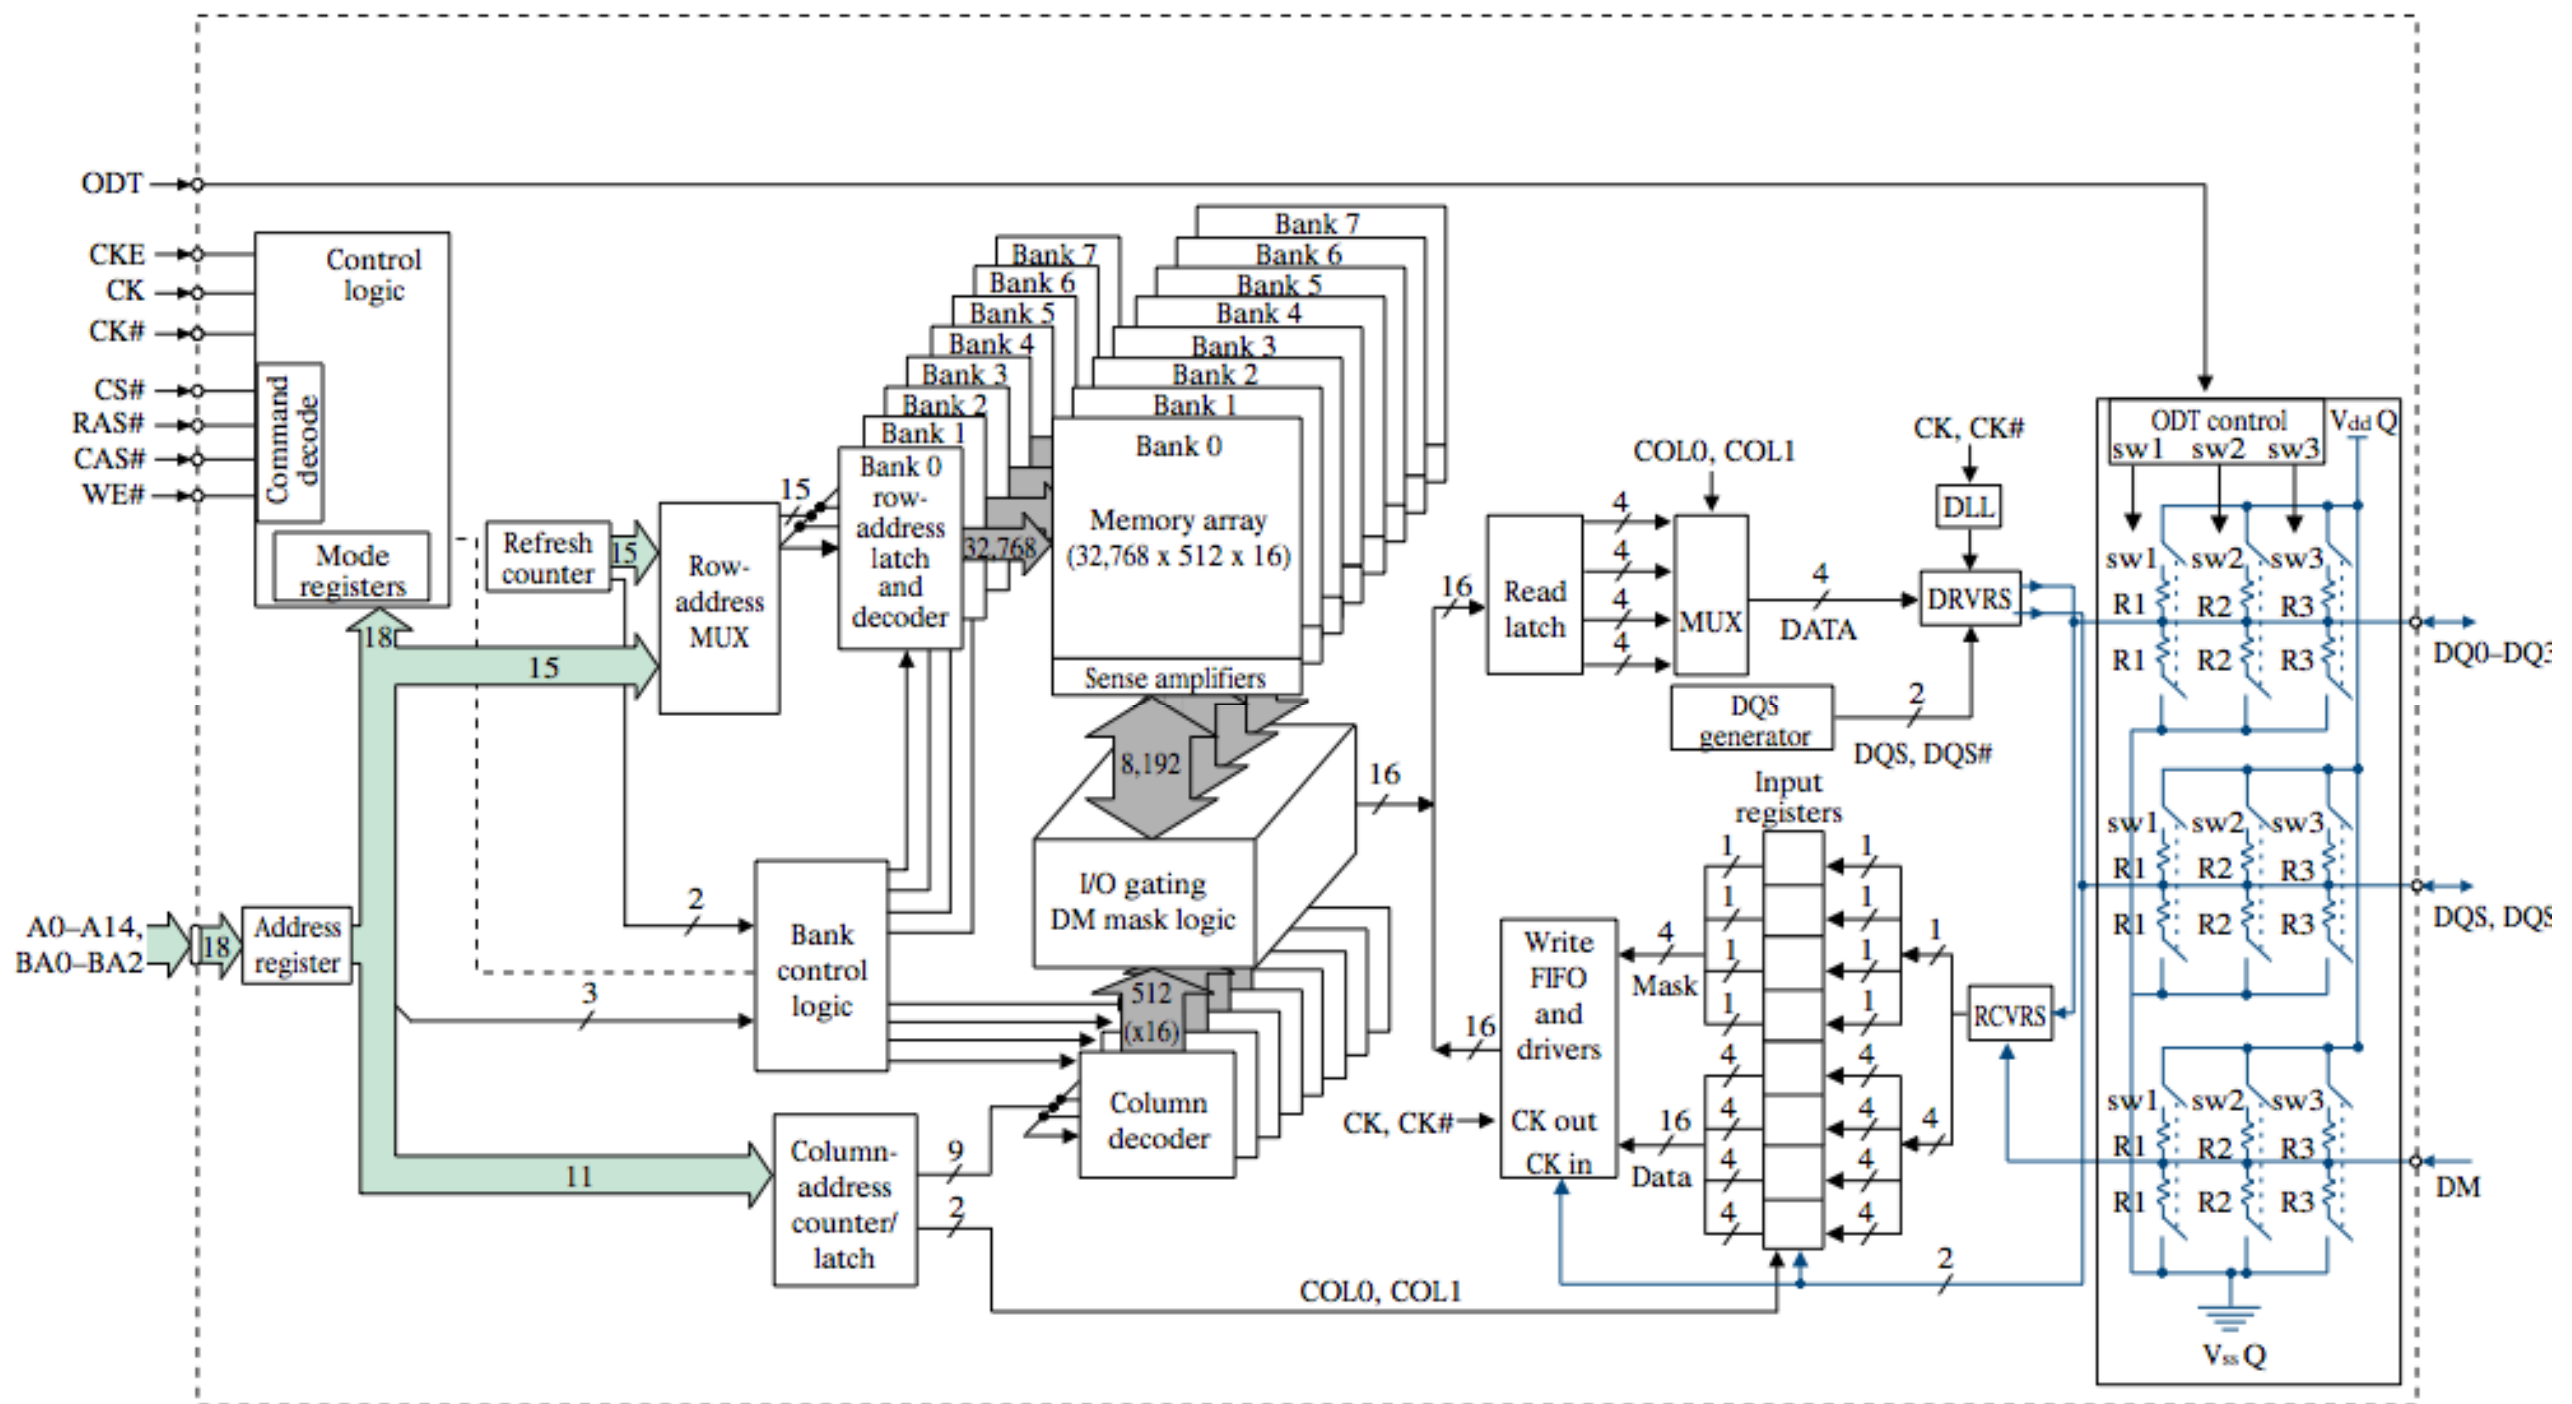
DRAM Devices

- There are many banks per die (16 at left)
 - Multiple pages can be open at once.
 - Can keep pages open longer
 - Parallelism
- Example
 - open bank 1, row 4
 - open bank 2, row 7
 - open bank 3, row 10
 - read bank 1, column 8
 - read bank 2, column 32
 - ...

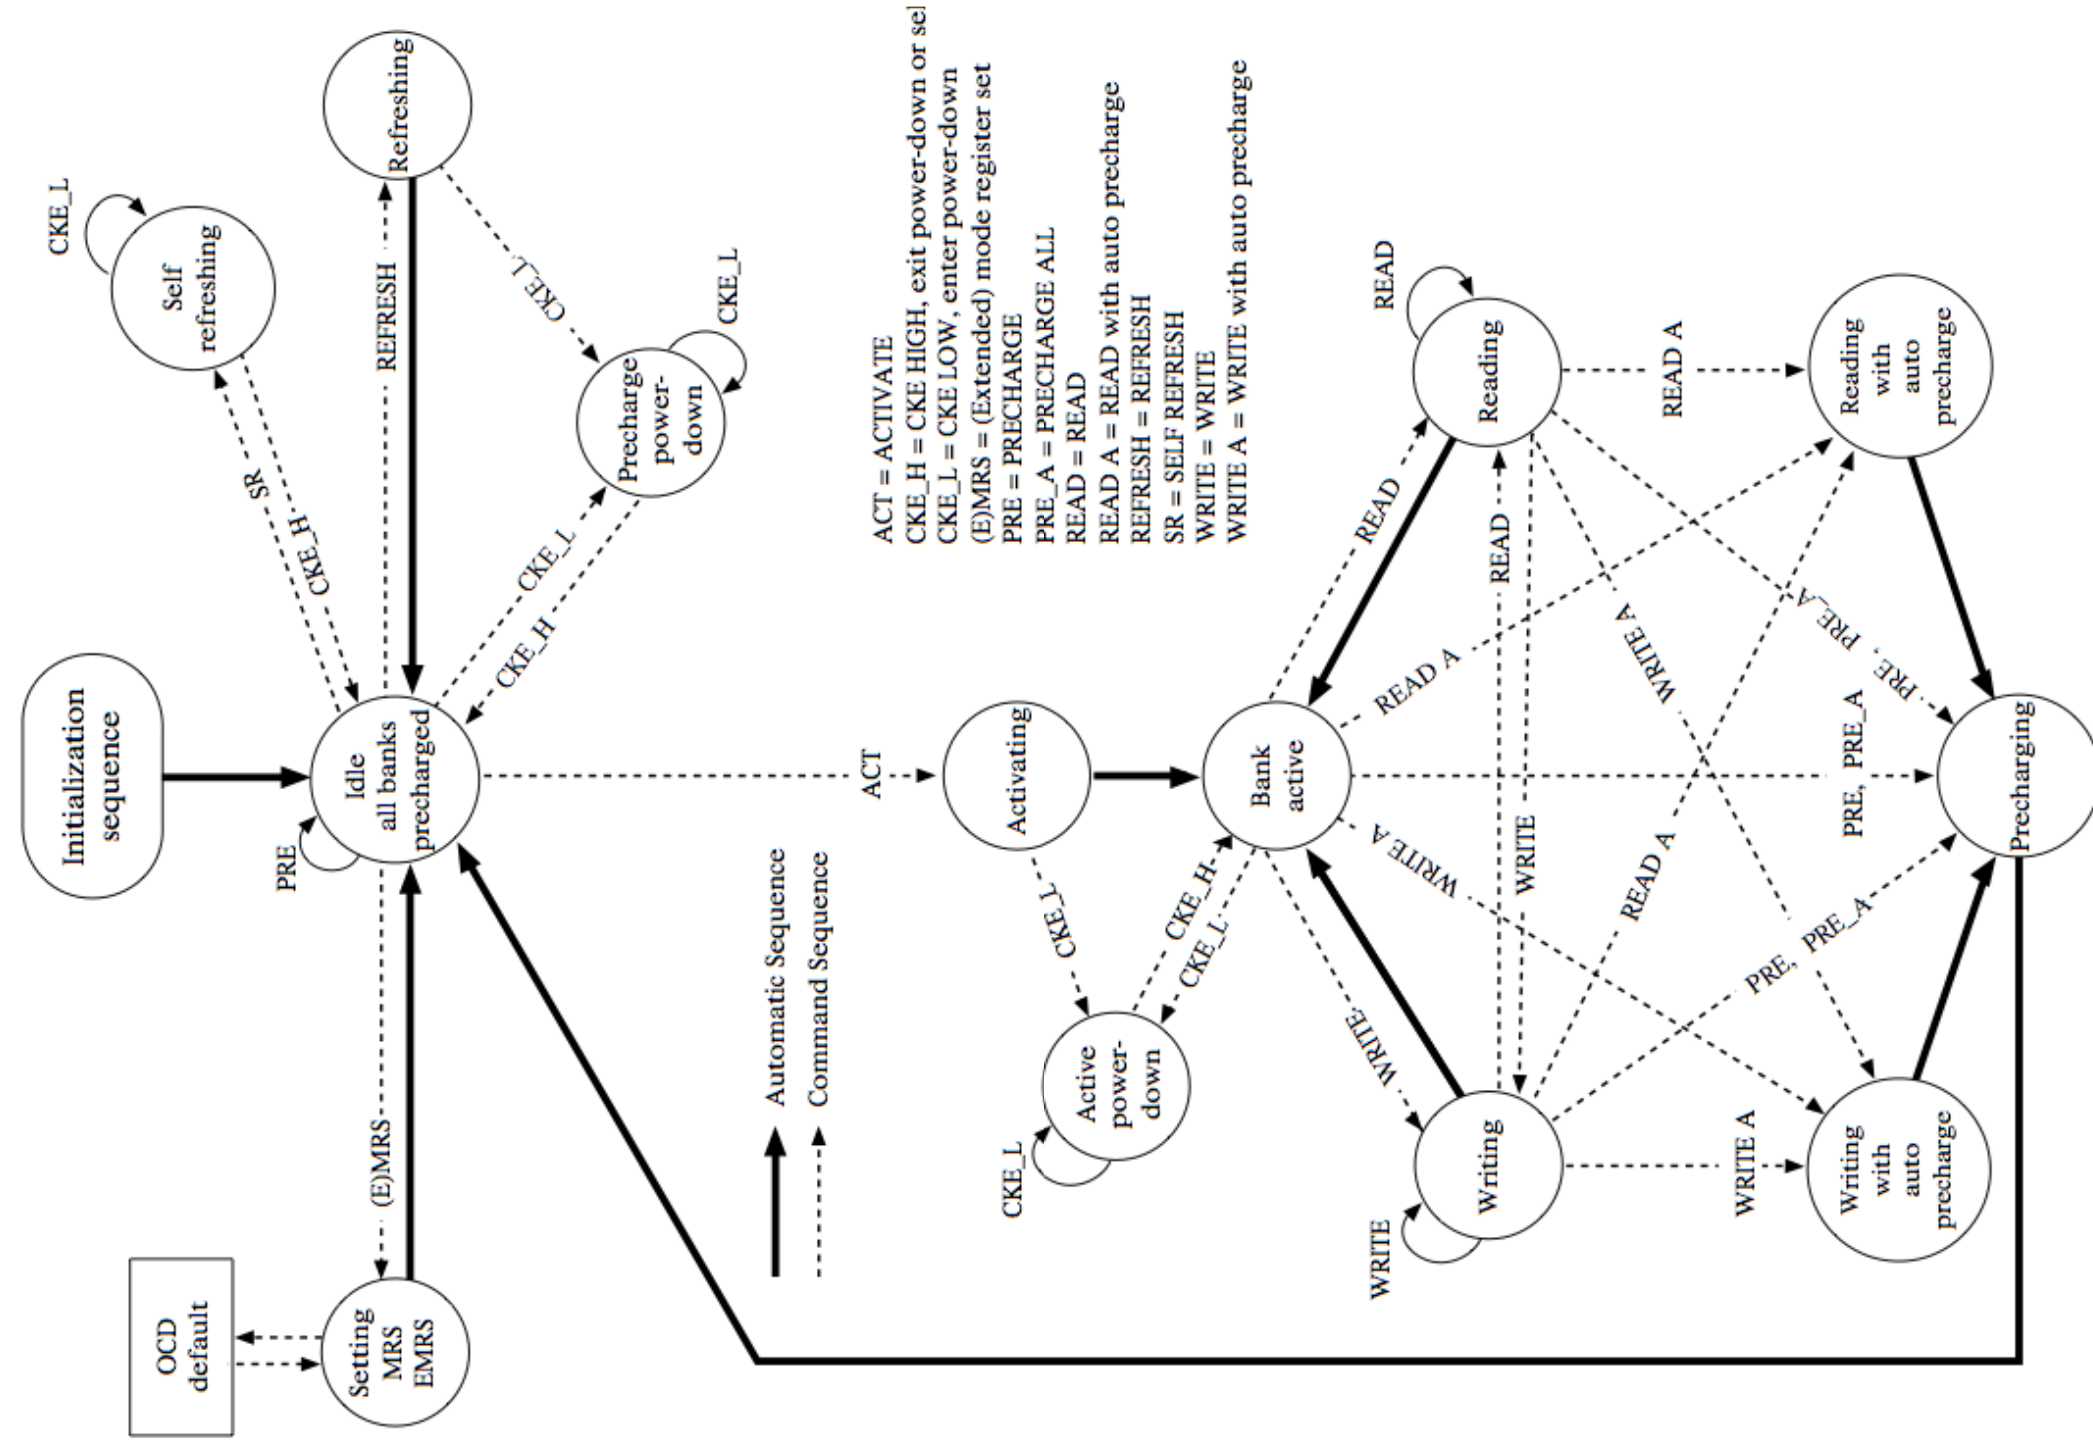


Micron 78nm 1Gb DDR3

DRAM: Micron MT47H512M4



DRAM: Micron MT47H512M4



DRAM Variants

- The basic DRAM technology has been wrapped in several different interfaces.
- SDRAM (synchronous)
- DDR SDRAM (double data-rate)
 - Data clocked on rising *and* falling edge of the clock.
- DDR2 -- faster, lower voltage DDR
- DDR3 -- even faster, even lower-voltage
- GDDR2-5 -- For graphics cards.

Current State-of-the-art: DDR3 SDRAM

- DIMM data path is 64bits (72 with ECC)
- Data rate: up to 1066Mhz DDR (2133Mhz effective)
- Bandwidth per DIMM GTNE: 16GB/s
 - *guaranteed not to exceed*
- Multiple DIMMs can attach to a bus
 - Reduces bandwidth/GB (a good idea?)



Each chip provides one 8-bit slice.

The chips are all synchronized and received the same commands

DRAM Scaling

- Long term need for performance has driven DRAM hard
 - complex interface.
 - High performance
 - High power.
- DRAM used to be the main driver for process scaling, now it's flash.
- Power is now a major concern.
- Scaling is expected to match CMOS tech scaling
- F² cell size will probably not decrease
- Historical foot note: Intel got its start as a DRAM company, but got out of it when it became a commodity.